Cultural Conflicts, Professional Identity, and Workforce Sorting: Evidence from U.S. Public School Employees

Roxanne Rahnama, Tyler Simko, and Elisa Wirsching*

August 26, 2026

PLEASE DO NOT DISTRIBUTE DRAFT

Abstract

How do cultural conflicts reshape the composition of the institutions they target? We argue that culturally conflictual policies threaten bureaucrats' public sector motivation and shape employment patterns through two channels: how a policy constrains an employee's discretion, and whether the ideological direction of the policy aligns with the employee's own values. We test our argument by collecting administrative roster data on 2.9 million public school employees across 12 states. We link each employee to a contentious policy adoption in their school district, and identify how each policy impacts their individual employment choices. Using a matched triple difference-in-differences design, we find that school district bans on the instruction of Critical Race Theory increase the probability that Democratic employees move to other districts, while Republican employees become more likely to remain in their current districts. Further, we find book ban adoptions drive profession exit among Democratic librarians. Interviews with current and former public school employees corroborate both channels, showing policy misalignment imposes costs even on educators whose own subject is not directly targeted. Our findings reveal how nationalized local political conflict can lead to durable partisan sorting in the bureaucracy.

*Rahnama is a Provostial Fellow and Lecturer at Stanford University (rrahnama@stanford.edu). Simko is an Assistant Professor at the University of Michigan (tsimko@umich.edu). Wirsching is an Assistant Professor at the London School of Economics and Political Science (e.m.wirsching@lse.ac.uk). We thank Luca Bellodi, Gwyneth McClendon, Cyrus Samii, Hye Young You, and Melissa Lyon. Kaleb Houle-Lawrence, Jayden Song, Pauline Huang, and Mia Coene provided excellent research assistance. This project was approved by Stanford University's institutional review board (IRB-FY2025-81590 and IRB-FY2026-87527) and generously supported by the Stanford Summer Research College and Princeton Data-Driven Social Science (DDSS) grant.

Introduction

A recurring question in the politics of democratic societies concerns whether public-facing institutions should maintain formal neutrality on contested social questions. In the United States, Americans across the political spectrum broadly favor political impartiality in principle (Yair and Sulitzeanu-Kenan, 2018; Torres, 2026). Yet, recent years have produced a sustained wave of legislation and executive action at national, state, and local levels that direct institutions like museums, national parks, schools, hospitals, and public and private universities to change their work and positions on questions related but not limited to race, gender, sexuality, and national identity – which scholars of cultural conflict identify as contests over whose values count as legitimate and are recognized in public life (Hunter, 1992; Wildavsky, 1987; Jacoby, 2014; Leege et al., 2009; Duttmann, 2000).

Educational institutions have long been sites of cultural contestation in the United States (Zimmerman, 2024). In its most recent iteration, state legislatures and school boards across the country have adopted binding restrictions and policies on curriculum, books, and classroom speech that declare specific forms of professional judgment impermissible, inviting national partisan conflict into public schools (Kogan, 2022; Torres, 2026). While existing work grapples with the drivers of these cultural conflicts (Lyon et al., 2026; Goncalves et al., 2024) and their consequences for student learning (Kogan, 2025), we know little about their consequences on the people working inside these politicized institutions. How do cultural conflicts reshape the workplace composition of the institutions they target?

In this paper, we study how culturally conflictual policies[1] — such as mandates on curriculum, books, and classroom speech — change employment patterns among U.S. public-school employees, the largest single group of "street-level bureaucrats" in the country (Lipsky, 1980; Maynard-Moody and Musheno, 2003). With over 5.4 million employees in K-12 schools, comprising

[1] Throughout the paper we use "culturally conflictual policies" and "cultural policies" interchangeably.

roughly 8 percent of the college-educated labor force (Kraft and Lyon, 2024), public school employees are frequently the direct targets of cultural policies (Lyon et al., 2026; Pawlewicz, 2020). Further, these restrictions directly impact professional discretion, given educators' roles as primary interpreters and implementers of policies in their classrooms (Weatherley and Lipsky, 1977; Hall and Hampden-Thompson, 2022).

We argue that cultural policies impose costs on public sector employees by threatening workers' ability to enact the professional identities they have developed around the exercise of their expertise. We argue two dimensions structure how employees experience these costs. First, *ideological alignment* captures whether the conflictual policy's change is congruent or opposed to the worker's values. Cultural policies constitute official public claims about whose professional values, judgment, and forms of knowledge are recognized as legitimate within the institution, which can shift the alignment between the institution's official position and an employee's personal ideology. Second, *role-based exposure* reflects how directly a policy targets the professional domain in which a specific worker exercises discretion. For example, book bans more directly target the discretion of librarians than athletic coaches. Under sufficient cost across these dimensions, workers reconsider their employment, potentially changing where they work or exiting the profession entirely (Hirschman, 1972).

Our analysis draws on a novel dataset of administrative public school employee roster data from 12 states covering approximately 2.9 million employees across school years 2017-18 through 2024-25. After acquiring each state's roster of employees via public records requests, we operationalize cultural policies in three ways. We identify school districts that (1) enacted bans against the instruction of Critical Race Theory (CRT), (2) adopted book bans that removed instructional material from their libraries, and (3) a broad set of controversial "culture war" incidents related to issues like gender equity, freedom of expression, and religion (from CATO). Then, we identify each employee's professional role (via our administrative data) and partisanship (by merging with voter files) to test our argument. Our classification is intended to capture our argument that professionals are differentially exposed to cultural policies based on their role, *and* that employees interpret cultural policies via their own political ideology.

We find that public school employees respond to cultural policies by geographically sorting between employers along partisan lines. In our main specification, we find that a school district's adoption of an anti-CRT policy exacerbates a political bifurcation of the public education workforce—politically misaligned employees become nearly twice as likely to leave their districts compared to baseline turnover rates, while aligned employees become *less likely* to move to other districts. These subgroup effects compound to a difference in the average treatment effect on the treated of 3.7 percentage points between politically misaligned and aligned employees. Further, consistent with our theory, we find these effects are strongest among employees whose work is most directly exposed to the content of the policy change, like English/Language Arts and History teachers.

We supplement our CRT analyses with directionally consistent, yet noisier, results for book bans and CATO-incidents as treatments. These alternative outcomes allow us to more closely examine sorting dynamics among more targeted treatments than CRT allows—for example, we find that book bans push politically misaligned librarians to completely exit their profession at much higher rates than aligned librarians. Our interviews with current and former employees support our core mechanisms and also point to barriers many teachers faced when converting the costs we theorize into employment changes, like vested pensions and the difficulties of starting new careers. Consistent with our qualitative evidence, our subgroups analyses reveal the largest employment effects are concentrated among teachers who faced the lowest barriers to moving and have less discretion to begin with (e.g. younger employees, who are more likely to be non-tenured). Further, we find that changes in turnover are largest in districts with higher proportions of low-income and non-white students. These districts already face significant turnover (Sutcher et al., 2019), with consequences for student learning (Sorensen and Ladd, 2020).

This paper makes three core contributions. First, we contribute to work on the politicization of education (Paglayan, 2026b; Pierskalla and Sacks, 2020), offering the first large-scale causal evidence on the consequences of educational politicization for the public school workforce. This adds to existing work that documents how cultural conflict in schools has influenced local governance, worsened student outcomes, and increased turnover among superintendents (Holman et al., 2025; Kogan, 2022; Peterson et al., 2025). Our results complement Paglayan et al. (2026), who shows that the most consistent feature of politicized education between regime types globally is the politicization of the teaching career (rather than curricula), including formal hiring and firing on political grounds. We make a within-democracy parallel from the perspective of educators' own responses to efforts of state control.

Second, we provide evidence for how cultural conflict influences public employees' decisions about where they work. Building on theoretical accounts of the importance of mission alignment between public workers and institutions (Besley and Ghatak, 2005; Forand et al., 2023; Gailmard and Patty, 2007; Gailmard and Gailmard, 2026), we study how changes in bureaucrats' non-material benefits through cultural policies affect bureaucratic selection. Existing empirical work on ideological misalignment between bureaucrats and political principals primarily focuses on routine administrative turnover (Richardson, 2019; Bolton et al., 2020; Spenkuch et al., 2023) or material politics (Wirsching, 2026a). Instead, we join an emerging body of work showing how symbolic politics and non-material workplace costs differentially shape selection among public employees (e.g. Brown et al., 2026).

Third, we demonstrate how cultural conflicts in schools exacerbate political sorting of the education workforce. Our results contribute to both the literatures on teacher turnover and sorting in education policy and economics—which has overwhelmingly focused on how workplace conditions like salaries, benefits, unionization, student composition, and accountability policies predict employment behavior (Biasi, 2021; Brunner et al., 2019; Barrett et al., 2021; Jackson, 2009; Jones and Hartney, 2017; Jeong and Luschei, 2019; Gjefsen and Gunnes, 2020; Nguyen et al., 2020; Moe, 2011, 2019)—and also residential political sorting—which has debated the extent to which Americans self-segregate into communities based on political ideology (e.g. Bishop and Cushing, 2009; Abrams and Fiorina, 2012; Mummolo and Nall, 2017; Martin and Webster, 2020; Brown and Enos, 2021). Our novel administrative data allows us to demonstrate that the education workforce is politically sorted above and beyond levels predicted by baseline local partisanship alone. While our measurements complement existing evidence about political sorting in the workplace (Chinoy and Koenen, 2024; Frake et al., 2026; Wirsching, 2026b), we extend this body of work by showing how the adoption of culturally contentious policies intensifies these patterns. Growing partisan divides on education issues (Houston, 2024) suggest that politicized policies consistent with our theory are likely to continue emerging in the future.

Theory and Conceptual Framework

Cultural conflicts are contests over whose values are given legitimacy and recognition in public life (Hunter, 1992; Gusfield, 1986; Jacoby, 2014; Wildavsky, 1987; Oldmixon, 2002) and arguments about "how we should live" (Leege et al., 2009, p.26). These conflicts can influence nominally neutral public-facing institutions by driving them to take official positions and establish policies on contested social questions. In doing so, cultural policies can invite ideological struggle into the workplace. For example, legislatures that enact binding policies on classroom content and speech in schools, condition funding on the (dis)use of certain language, or bar medical providers from offering gender-affirming care or discussing abortion referrals in hospitals are making official public claims about what professional judgment and knowledge count as valid within the institution. We argue transformation can impose costs on workers through two channels: by shifting the alignment between the institution's official position and workers' professional mission, and by reducing the professional discretion through which workers enact their expertise. These costs fall unevenly across the workforce and may lead some to reconsider their employment.

Professional Identity and Why it is at Stake

Understanding why these official claims impose costs requires starting with professional identity. Identity refers to an actor's self-definition: the way an individual answers the question "who am I?" (Ashforth and Schinoff, 2016; Ashforth, 2000; Stryker and Burke, 2000). For professionals, this self-definition is particularly organized around what they do. Because society grants professionals prestige and autonomy in recognition of their specialized knowledge and skills (Pratt et al., 2006; Freidson, 2001; Larson, 1977), workers in skilled occupations come to see the exercise of professional expertise as an expression of who they are, internalizing the values and standards of their fields as constitutive of their self-concept (Van Maanen and Barley, 1982; Ibarra, 1999). These attachments develop through on-the-job socialization that shapes role expectations and personal sensemaking, i.e., the process through which individuals construct a coherent sense of self by making meaning of their daily work (Ashforth and Schinoff, 2016; Maitlis and Christianson, 2014; Pratt, 2000).

Institutional contexts can disrupt the conditions under which workers enact their professional identity, by reshaping whether and how the institutional setting recognizes and accommodates their professional self (Ashforth and Schinoff, 2016). Sensebreaking and sensegiving describe how institutions rupture the meanings workers have organized around their roles: when organizational change, role shifts, or external pressure alter these conditions, workers reassess whether they can sustain their professional identity in their current institutional context (Pratt, 2000; DiBenigno, 2022). When the tasks performed are consistent with one's values, expertise and sense of purpose, the alignment between identity and work produces intrinsic rewards that material compensation cannot substitute (Rosso et al., 2010; Pratt et al., 2006). Workplace disruption can create tension between this self-conception and what their new institutional context allows employees to do (Brown et al., 2026). We study how cultural policy changes can disrupt this professional identity.

Public Service Motivation and its Double Edge

We argue that workers in public-service institutions should be sensitive to these institutional claims. Street-level bureaucrats—workers who deliver public services through direct contact with citizens (Lipsky, 1980; Maynard-Moody and Musheno, 2003)—derive substantial public service motivation (PSM) from the experience of their work as purposeful (Perry and Wise, 1990; Gailmard and Patty, 2007; Ashraf and Bandiera, 2018). PSM is defined as an individual's predisposition to respond to motives that are grounded mostly in public institutions, i.e., a general orientation toward civic duty and meaningful public contribution that draws individuals into public service and sustains their commitment to it (Perry and Wise, 1990). To fix ideas, consider a worker's simple utility from working in the public sector that combines wage and material benefits (w)—the primary focus of existing research on public sector employment decisions (Biasi, 2021; Moe, 2011)—with non-material mission value (v):

$$u^P = (1 - p) \times w + p \times v \tag{1}$$

Since PSM (p) scales the value workers attach to non-material mission benefits, those with stronger public service identification feel the rewards of alignment and the costs of mission violation more acutely than their lower-PSM counterparts. Workers remain when this utility exceeds the value of the next-best alternative and exit or reconfigure their employment when it does not (Hirschman, 1972).

Critically, high PSM is not only what draws workers into public service but also what makes them vulnerable when institutional claims move against their values. As Bunderson and Thompson (2009) and Schabram and Maitlis (2017) argue, callings are double-edged: they bind workers to their work precisely because the work is experienced as a constitutive expression of who they are, but this same fusion of identity and vocation makes mission violations more costly. Value incongruence and constraints on professional practice depress job satisfaction more sharply for high-PSM workers (Tummers, 2012; Hayes and Stazyk, 2019), and the attraction-selection-attrition process through which mission-aligned workers disproportionately sort into public institutions (Schneider, 1987) means that those with the strongest organizational commitment are the most exposed to costs when institutional missions shift (Gailmard and Gailmard, 2026).

Two Dimensions of Non-Material Benefits: Alignment and Discretion

We argue that the non-material mission value that PSM activates operates through two analytically distinct dimensions:

$$v(a,d) \;=\; \underbrace{\alpha(a)\,a}_{\text{symbolic}} \;-\; \underbrace{\beta\,e\,(1-d)\,g(a)}_{\text{constrained practice}} \tag{2}$$

The first dimension is *alignment* a: the degree of congruence between a worker's core values and her organization's official mission and position. Meaning arises from perceived congruence between workers' values and those of their organizations (Rosso et al., 2010; Besharov, 2008; Pratt, 2000), such that this alignment is itself a source of non-material compensation (Besley and Ghatak, 2005; Carpenter and Gong, 2016; Khan, 2025; Vandenabeele, 2008; Wilson, 1989). Workers who perceive their organization's mission as reflecting their own values derive positive meaning from institutional membership (Hayes and Stazyk, 2019). This alignment value is independent from the domain a public employee works in and whether an institution's official policy regulates their individual behavior. However, this relationship may be asymmetric. Indeed, Rosso et al. (2010, p.104) discuss a form of loss aversion, where "organizations that form psychological contracts with employees on the basis of mission or ideology may face a stronger set of negative reactions if they are perceived to be violating these ideals" than positive reactions that mission congruence generates.

The second dimension is *discretion* d: the extent of freedom a worker can exercise in her professional role (Lipsky, 1980; Honig, 2024). Professions that allow for higher levels of autonomy, skill variety, and task significance produce greater experienced meaningfulness (Hackman and Lawler, 1971; Rosso et al., 2010). More importantly, discretion is a primary means through which workers bridge the gap between what they believe is professionally right and what their organization requires of them, $g(a) \geq 0$. This gap is wider the more her values diverge from the institution's official position, and vanishing when the two coincide. Discretion d is the share of that gap she is free to close, so $(1 - d)\, g(a)$ is the residual distance between what she believes is right and what she is permitted to do.[2] Workers proactively adapt the boundaries of their tasks to bring mandated practice closer to their professional ideal (Wrzesniewski and Dutton, 2001). For policy-motivated public sector workers, discretion is a form of compensation that is uniquely available in public service, as it allows them to advance causes they consider important and bring policy outcomes closer to their own ideal points (Gailmard and Patty, 2007). Additionally, discretion only matters for workers whose professional domain is affected by the institution's official position, as captured by $e \in \{0, 1\}$. For a worker whose practice is not directly implicated by a policy (e.g., a mathematics teacher, where the contested question is the treatment of race in the history curriculum), the second term is inactive. Instead, her mission value is driven by alignment alone.

How Culturally Conflictual Policies Disrupt Both Dimensions

Cultural policies threaten alignment and discretion in several ways.[3] First, cultural policies shift the alignment dimension for all workers in the targeted institution.[4] Because professional

[2]We omit any independent value of discretion for simplicity. Including one would add a constant cost for all directly exposed workers, but would leave the interaction between exposure and alignment — the quantity our later research design identifies — unchanged.

[3]The policies we study restrict what workers may do and say in their professional roles rather than targeting material resources, unlike policies that operate through material deprivation — for instance President Trump's FY 2027 proposal to cut "woke programs" (Broadwater and Bender, 2026). Both embed a symbolic legitimacy claim but may impose different behavioral constraints. How the consequences of practice-restricting cultural policies compare to material deprivation-policies, and how both benchmark against non-ideological constraints, is an important area for future research.

[4]A scope condition of our argument is enactment of the policy: a cultural conflict imposes institutional costs when it produces an enacted policy, rather than remaining a contested but unresolved public struggle. Enactment is what transforms an ideological conflict into an official institutional claim.

standing rests in part on the institutional recognition of specialized expertise as worthy of deference and autonomy (Larson, 1977; Freidson, 2001), this official claim invokes a form of sensebreaking for workers whose values the enacted position repudiates (Ashforth and Schinoff, 2016; Pratt, 2000). It ruptures the meaning they had built around their roles and withdraws the legitimation on which professional self-concept depends, because they are no longer treated as professionals whose expertise warrants recognition. For workers whose values it affirms, the same official act runs in the opposite direction, providing institutional recognition to a professional self-understanding that had previously gone unendorsed.

In terms of equation (2), enactment therefore shifts a for every worker in the institution. The shift is signed — negative for those the official position moves away from, positive for those it moves toward. Crucially, this shift reaches workers whether or not their own professional domain is regulated. The institution has declared something about what counts as legitimate within it, and that declaration reaches the entire workforce.

Second, restrictions reduce discretion (d). They do so directly, by prohibiting content and practices a worker would otherwise have been free to choose, and indirectly, through their enforcement machinery, such as complaint procedures and exposure to formal sanction, which induces self-censorship beyond the restriction itself. A worker required to withhold guidance she regards as professionally essential, or to implement practices she considers harmful to those she serves, may in turn experience moral injury (the psychological cost of being required to act against one's own values), policy alienation, and the vigilance costs of continuously monitoring her own practice for sanctioning risk (Čartolovni et al., 2021; Zacka, 2017; Tummers, 2012).[5]

Unlike the alignment shift, this second disruption is role-dependent. It reaches only those workers whose own professional domain is regulated by the policy. And because $(1 - d)$ multiplies $g(a)$

[5]Surely, workers arrive at a policy already carrying other strains to their workplace environment that affect their non-material benefits from public sector employment, including workload policies or the relationship with and behavior of clients (e.g., a teacher's students). Our claim is about the marginal effect of the policy, not the worker's total satisfaction or distress, and as long as these other factors are not highly correlated with the policies we care about, they pose no challenges to test our theory.

in equation (2), the cost of losing discretion is proportional to the gap it would have been used to close. The same restriction is expensive for a worker whose professional ideal departs sharply from what the institution now mandates, but a worker whose ideals are now closer to the mandated practice benefits from that change: she is released from a requirement that had previously created a gap in her practice. Exposure and discretion thus amplify the alignment shift in both directions, compounding the loss for the misaligned and the gain for the aligned. Conversely, just as the value of discretion depends on alignment, the value of alignment depends on discretion. Discretion insulates a worker from her institution's official position; a worker with little latitude absorbs the full gap, so a shift in that position bears on her more directly. Alignment and discretion are thus substitutes: each matters more where the other is scarce.[6]

Who Bears the Costs and Who Reaps the Benefits

The two disruptions, therefore, fall on different workers, and their intersection determines who bears what. Table 1 illustrates the four cases that our argument predicts.

We predict that workers who are both ideologically misaligned and directly exposed to cultural policies face the strongest exit pressure. These employees experience restrictions on their own professional discretion with which they ideologically disagree, producing the highest predicted probability of exit. Among workers who are ideologically misaligned but not directly exposed, only the alignment dimension is disrupted; their own practice is untouched, and exit pressure is correspondingly more moderate.

For ideologically aligned workers, by contrast, our framework predicts a *gain*, which is largest for those in directly targeted roles. The policy both affirms their values and reduces the identity gap the worker needs to bridge with discretion.[7] This directional prediction is what separates

[6]Formally, $\partial v/\partial d = \beta e\, g(a)$ and $\partial v/\partial a = \alpha - \beta e(1-d)g'(a)$. The interactions operate only where $e = 1$; for workers whose domain the institution's position does not govern, mission value is purely symbolic and depends on alignment alone.

[7]The aligned, exposed worker also faces reduced discretion, but less discretion is needed to bridge an already closed identity gap g(a).

Table 1. Predicted employment change by role-based exposure and ideological alignment

	Ideologically aligned	**Ideologically misaligned**
Non-exposed	The institution's position moves toward her values; her own practice is unaffected. ***Net gain; retention.***	The institution's position moves away from her values; her own practice is unaffected. ***Moderate exit pressure.***
Directly exposed	The institution's position moves toward her values *and* mandated practice moves with it, reducing the identity gap needed to bridge ***Largest gain; strongest retention.***	The institution's position moves away from her values, *and* new mandates reduce employee discretion needed to bridge a larger identity gap. ***Highest exit pressure.***

our account from explanations that operate through disruption, red tape, or sanctioning risk alone (Honig, 2024; Brown et al., 2026). Those mechanisms can generate higher exit among misaligned workers, but do not predict *improved* retention among aligned workers in the same institutions at the same moment.

Strictly speaking, these predictions hold ceteris paribus, with workers' outside options held fixed. Street-level bureaucrats turn over only once these costs grow large enough that the value of the public sector position falls below that of the next-best alternative. We assume that—and our later research design helps assure—-opportunity costs and outside options are approximately balanced across the comparisons we draw, particularly for employees with similar characteristics, including subjects, grades, and states, in treated and control school districts.

Case and Context

Public School Employees as the Case

We apply this argument in the case of public-school employees for several reasons. First, public-school educators are high-PSM workers who often describe their profession as a calling (Johnson, 2004), deriving substantial motivation from work experienced as meaningful and consistent with their mission values. Extensive research has shown that salary is rarely the key driver of educators' career choices (McNeil, 2015; Kraft and Lyon, 2024). Professional discretion over instruction is central to how educators fulfill this public mission (Mehta, 2013; Lopes and Oliveira, 2020; Kraft and Lyon, 2024; Dunn, 2018; Worth and Van den Brande, 2020) and to how they enact their professional identity, making principled choices, translating curriculum into practice, and responding to the needs of students on a daily basis (Dunn, 2018; Pignatelli, 1993; Anderson, 2010).[8]

Second, cultural policies in education affect public school employees' professional identity through both channels of our theory – making official institutional claims about whose professional judgment counts as legitimate and restricting the discretion through which employees in targeted roles enact that identity in practice. Adopted local policies like book bans are institutional claims that educators' professional judgment in their own domain is insufficient. As Renkl (2022) claims, an "unspoken belief underlying such ideological policing is that...teachers don't deserve to be regarded as the skilled professionals they are." That such policies have been adopted at a wide scale and pace (Woo et al., 2024) makes the current period a valuable moment to study the workforce consequences.

Third, public school employees face unusually high institutional costs from changing employers. Moving across districts, for example, can typically lead to forfeiting tenure protections and resets pension accrual under defined-benefit plans that tie retirement benefits to years of service within a single system; departure before reaching threshold years of service can incur a penalty (Koedel

[8] Dunn (2018) draws insights from teachers' public resignation letters, where common themes that emerged include a "mismatch between their beliefs and the reality of teaching in today's educational climate, and a lack of trust and respect for their profession, and a lack of control over their working conditions" (p.1).

et al., 2013; National Education Association, 2023). According to the National Center for Education Statistics (NCES), roughly 8 percent move schools each year, with a majority being voluntary transfers within districts (National Center for Education Statistics, 2024). Professional exit (from a district or the profession) is therefore a difficult outcome to move and, as a result, a difficult test of our argument,[9] especially for targeted educators in non-STEM fields who have fewer exit options (Nguyen et al., 2020).

Cultural Conflicts in Education

Cultural conflict over what American children learn in school is not new. Curricula are not neutral, often embodying values and beliefs that are constantly contested (Ravitch, 1990). Primary school curricula have been sites of intense political contestation since at least the nineteenth century (Paglayan, 2026a), and the recurring themes—race, gender, sexuality, religion, and national identity—have produced a long historical arc with recognizable patterns (Lyon et al., 2026). White Southern elites systematically controlled textbooks to promote Lost Cause narratives during Jim Crow (Rahnama, 2025); Cold War-era conflicts targeted "subversive" classroom content; evolution-creationism battles recurred across decades (Berkman and Plutzer, 2010); multicultural education reforms of the 1960s through 1990s were characterized by the "New Right" as un-American (Zimmerman, 2024); and the 1970s "Save Our Children" campaign sought to ban gay teachers from public schools (Rofes, 2005). Reforms throughout the history of the American common school have characterized educators as both the problem and the solution to perceived shortcomings of public education, repeatedly producing efforts to raise instructional quality by controlling their practices through top-down management and standardization that diminish regard for their expertise (Mehta, 2013; Pawlewicz, 2020).

The contemporary wave represents one of the largest and most geographically extensive mobiliza-

[9] Charter schools and private schools represent an intermediate outside option, where they need to follow state education mandates but are typically shielded from local school board rules. In practice, however, 96% of public school teachers who have changed schools moved to another public school (National Center for Education Statistics, 2024).

tions of this kind in recent American history. For example, between 2020 and 2021, 894 school districts that cover 35 percent of all K-12 students were impacted by local actions against critical race theory in classrooms (Pollock et al., 2022), accompanied by extensive related actions around "parental rights" in local school boards by groups like Moms for Liberty (Sinha et al., 2023; Goldstein, 2024). These various policies affect approximately 1.3 million public school employees and have accompanied the removal of nearly 16,000 books from school libraries nationwide (PEN America, 2024). Many of these policies concentrate on content derived from Trump's 2020 executive order on "Combating Race and Sex Stereotyping" — prohibiting instruction on "divisive concepts" such as systemic racism, sex-based privilege,and LGBTQ-related classroom content.

Main Hypotheses

Our argument above presents four core testable hypotheses:

H1 Ideologically misaligned public-school employees in districts that enact cultural policies are more likely to change their employment than ideologically aligned employees.

H1A Ideologically misaligned employees in districts that enact cultural policies are *more* likely to change their employment than comparable misaligned employees in districts that do not.

H1B Ideologically aligned employees in districts that enact cultural policies are *less* likely to change their employment than comparable aligned employees in districts that do not.

H2 The probability of employment change following enacted cultural policies is highest among employees who are both ideologically misaligned and in directly exposed roles, intermediate among those who are misaligned but not directly exposed, and lowest among ideologically aligned employees.

Data

We collect, clean, and process several different data sources to test our argument. Figure 1 describes our data processing pipeline. For both our treatment (instances of cultural policies) and outcome (teacher turnover), we collect large-scale data across the country and harmonize data across sources.

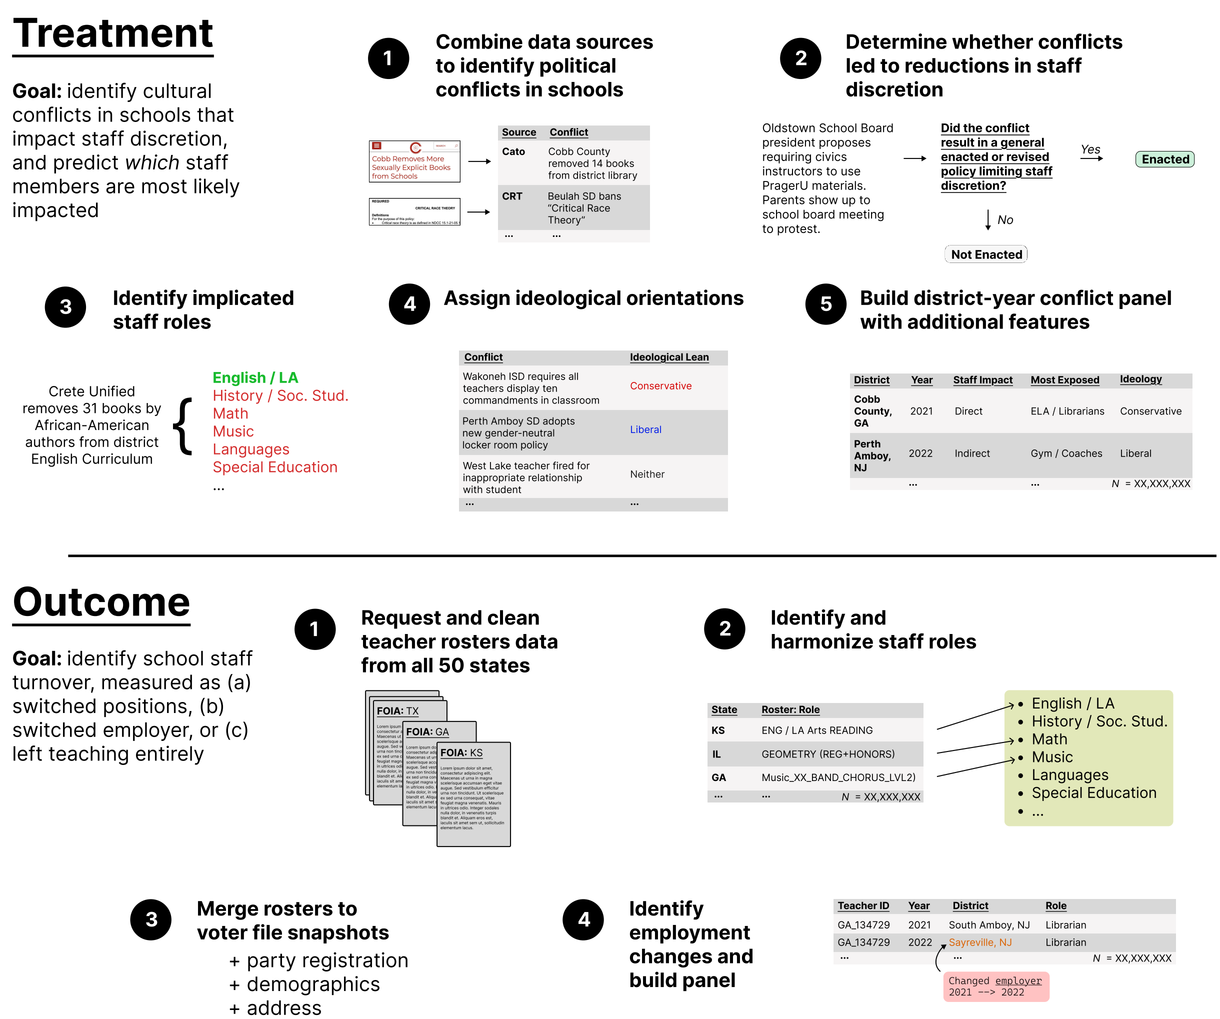

Figure 1. Data processing pipeline: The figure illustrates how we clean and process the data used to generate our treatment (top) and outcome (bottom) measures. Each step is described in detail in the text.

For treatment, our primary measurement goals are to (1) identify cultural policies in schools,

(2) identify the employees most likely to be implicated by these policy changes using a novel role categorization scheme, (3) assign ideological orientations to these policies, and (4) build a district-level panel of cultural policies across the United States. For outcomes, we (1) collect teacher and school staff rosters from states across the country, (2) use the same role categorization scheme above to assign every public school employee to a specific role (e.g. English/Language Arts, History/Social Studies, etc.), (3) merge employees to voter file snapshots for partisan and demographic information, and (4) build a panel of employment changes over time. In our research design, we then combine these two sources to examine how these policies influence employee turnover. Below, we describe each of these steps in more detail.

Treatment: Measuring Culturally Conflictual Policies

The first step of our pipeline in Figure 1 is to identify cultural policies in US school districts that restrict the discretion of teachers and other school staff during the course of their work. We also seek to identify how educational staff are differentially exposed to given conflicts.

We analyze three sources of cultural policies, which together allow us to examine turnover responses across different treatment types and patterns of role-based exposure. For each source, we build a district-year panel of restrictions and classify the ideological direction of each policy.

Critical Race Theory: Our primary treatment is anti-CRT restrictions, taken from UCLA's CRT Forward database (Alexander et al., 2024), which records state- and district-level bills restricting instruction of Critical Race Theory.[10] These policies are among the most geographically widespread and politically salient curricular conflicts in our sample period, and they are well-suited to our framework for three reasons. First, they occur frequently across many

[10]CRT is an academic framework, developed primarily in legal scholarship, that examines how racism is embedded in legal systems and social structures. As Bissell (2023, p.299) notes, "virtually no K-12 public school actually teaches CRT—but some may place an increased focus on diversity and multiculturalism, add important historical events like the Tulsa race massacre to their curricula, or seek to better support students of color. Proponents of anti-CRT bills have successfully conflated these activities with CRT." This conflation is argued to be, in part, deliberate: Christopher Rufo, widely credited as an architect of the anti-CRT legislative campaign, wrote in 2021 that the goal was to "freeze the brand of CRT and use it to encompass all of the various cultural insanities" (p.228).

districts over time, providing the geographic and temporal variation our design requires. CRT Forward captures several hundred such policies across 45 states as of December 2024.[11] Second, their ideological direction is clear and manually verifiable: we confirm case-by-case that each incident reflects a conservative restriction on curricular content (such as a prohibition on teaching CRT-related material) rather than a liberal mandate to include diverse perspectives; Appendix A describes this process in detail. Third, and most directly relevant to the theory, anti-CRT policies operate through the two channels we identify. They constitute official institutional claims about whose professional judgment is recognized as legitimate within the school, and they place binding oversight on what educators may present or say in the classroom, thereby restricting the discretion of targeted workers (particularly English/LA, history/social sciences, and librarians).[12] For example, in 2022 Cobb County School District in Georgia enacted a "Parents' Bill of Rights," allowing parents to "object to instructional materials intended for use in his/her child's classroom or recommended by his/her child's teacher."[13] We restrict attention to bills targeting K-12 instruction, dropping those that exclusively target other institutions such as higher education.

Book Bans: Our first secondary treatment comes from PEN America's Book Bans database (hereafter PEN Books). These data track book-level instances where instructional use of a book is restricted or diminished, with information on location, challenge and removal timing, ban status, and content (PEN America, 2024),[14] allowing us to evaluate whether a school district had any book bans or not. Book bans share the features that make anti-CRT policies useful—a verifiable conservative direction and a direct restriction on professional practice—but they are

[11] The methodology combines systematic searches of legal databases for state and federal legislation with screening of over 4,000 newspapers and 12,000 media sources to identify local-level measures not captured in existing databases. The database was updated monthly until the close of the project in late 2024.

[12] From a legal perspective, the government can set the curriculum to regulate the 'what' of teaching, but 'how' to teach is largely left to teachers as certified professionals. Yet, as Bissell (2023) discusses: "anti-CRT [policies] represent a sharp departure from this approach, reaching beyond curriculum to regulate the very words that come out of a teacher's mouth" (p.211).

[13] See the enacted Cobb County Parents' Bill of Rights here.

[14] The methodology captures various forms of censorship beyond outright removal, including books requiring parental permission, grade-level restrictions, or temporary removal for "review."

less geographically widespread, largely followed after CRT bans, and implicate a narrower set of roles,[15] with librarians being the most directly exposed group. This allows us to test the exposure channel with a more sharply defined treatment population. In 2022, for instance, the Washington Township Board of Education in New Jersey voted 8-1 to remove Toni Morrison's *The Bluest Eye* from English courses following a parent complaint over sexually explicit content.[16]

Additional Culturally Conflictual Policies: Our second secondary treatment is the Cato Institute's School Battle Map (hereafter Cato), compiled from 2001-2024 to identify school conflicts relating to basic rights, moral values, or individual identities (Kogan, 2022). Cato is noisier than our other sources: ideological direction is harder to verify, and not all incidents involve direct restrictions on professional practice. In exchange, it offers far wider geographic and temporal coverage and captures cultural policies the other sources overlook. In 2016, for example, the Youngstown School District in Ohio adopted a policy eliminating "any reference to intelligent design, creationism, or any like concepts" from the science curriculum. It therefore allows us to examine our theory's application to a broader set of cultural policies exposing a wider range of bureaucrats. As with the other two sources, we classify the ideological lean of each policy (pro-conservative, pro-liberal, or neither), here using an LLM (Claude Sonnet 4.6) applied to the event description.[17]

Identifying role exposure and ideological direction of policy changes

After we identify culturally conflictual policies in schools, testing our theory requires measures of both *role exposure*—identifying which roles are likely to be most impacted by specific changes—and *ideological orientation*—the perceived ideological direction of the final policy.

[15] Although CRT- and book bans share overlapping exposed groups, CRT bans also implicate history/social studies educators, while book bans do not.

[16] See more information on the Washington Township book ban here. The most commonly banned books in the PEN data from 2021 are *Gender Queer: A Memoir* by Maia Kobabe, *All Boys Aren't Blue* by George M. Johnson, *Out of Darkness* by Ashley Hope Pérez, *The Bluest Eye* by Toni Morrison, and *Lawn Boy* by Jonathan Evison.

[17] A research assistant blindly hand-coded 400 randomly selected cases out of 4345 to evaluate human-coding against Claude's classification (see Appendix A.2 for details on precision and recall on the coding dimensions).

In Steps 3 and 4 of the treatment section of Figure 1, we use a combination of manual hand-coding and large-language models (LLMs) to code role exposure and ideological orientation. This scheme is described in more detail in Appendix B. First, we inductively iterate between our public school employee roster and cultural policy data to develop a categorical scheme of the most common staff roles in U.S. school districts (e.g. English/Language Arts, History/Social Studies, Music, etc.).

Then to measure role exposure, we use a combination of manual coding and LLMs classify each policy in our dataset according to which positions (e.g. teacher, principal, counselor, etc.) and subjects (e.g. English/Language Arts, Science, Arts, etc.) are most likely to be implicated. Appendix Table A2 shows several examples of cultural policies and the roles we classify as most directly exposed. In the next section, we describe how we take a similar approach for classifying all 2.9 million school employees in our dataset. We do not argue—and our later analyses do not assume—that non-exposed groups are unaffected by district conflicts. For example, gym teachers may perceive arguments over book bans as generally disruptive to their workplace environment even if their classrooms are not directly impacted. However, the purpose of this predictive step is to identify the group most clearly and directly impacted by each change, if any.

For ideological orientation we identify whether cultural policies were in a conservative, liberal, or neither direction. For example, a book ban that restricts LGBTQ+ reading material is classified as pro-Conservative, while a policy change allowing gender-nonconforming students to use the locker room of their gender identity is classified as pro-Liberal. Examples of each variable are shown in Appendix Table A2.

Finally, in Step 5 of Figure 1 we aggregate these treatments to the district-year level.

Outcomes: Measuring Employee Turnover

To measure employment changes of teachers and other public school staff, we submitted public records requests to state educational authorities across all US states. We collected extensive administrative data on all employees of public school districts for the school years 2017/18 to 2024/25 which allow us to trace employees' careers and labor market dynamics. Specifically, for each US state, we requested data on employees' names, race, sex, birth dates, school and district of employment, educational background, subject assignments and titles, employment histories, hire and exit dates, and salaries. These data are structured as repeated, annual cross-sections of the states' public school workforce. Figure 2 shows the distribution of data availability by state. 29 states have provided detailed records, while 12 states have denied the requests and 9 states remain pending. In this iteration of the project, we focus on states with noticeable variation in educational culture wars in recent years, including Florida, Texas, Iowa, Idaho, Georgia, Kansas, Kentucky, Connecticut, Illinois, Missouri, Wisconsin, and Pennsylvania.

Identify and harmonize staff roles: We then map our scheme classifying the cultural policies along exposure to the public school employee rosters. As detailed in Appendix B, we categorize all 2.9 million public school staff in our sample into one of 10 non-mutually exclusive role categories (teachers, substitute teachers, librarians, counselors, coaches and extra-curricular staff, special-education staff, support staff, superintendents, administrative staff, and other positions) and 8 broad subject categories that also determines which culture-war front the staff member is exposed to (English language classes; history and social studies; health and physical education; science classes; arts and music; mathematics; world language classes; other subjects). This approach is necessary because school districts vary dramatically in how they organize and refer to their staff—for example, administrative records in Georgia refer to librarians with varied titles like `Librarian`, `Media Specialist` and `Media Secretary/Clerk`, while Kentucky calls them `Media Librarian`. One important caveat of this exercise is that state administrative records vary considerably in their level of detail and structure. To ensure most consistent classifications

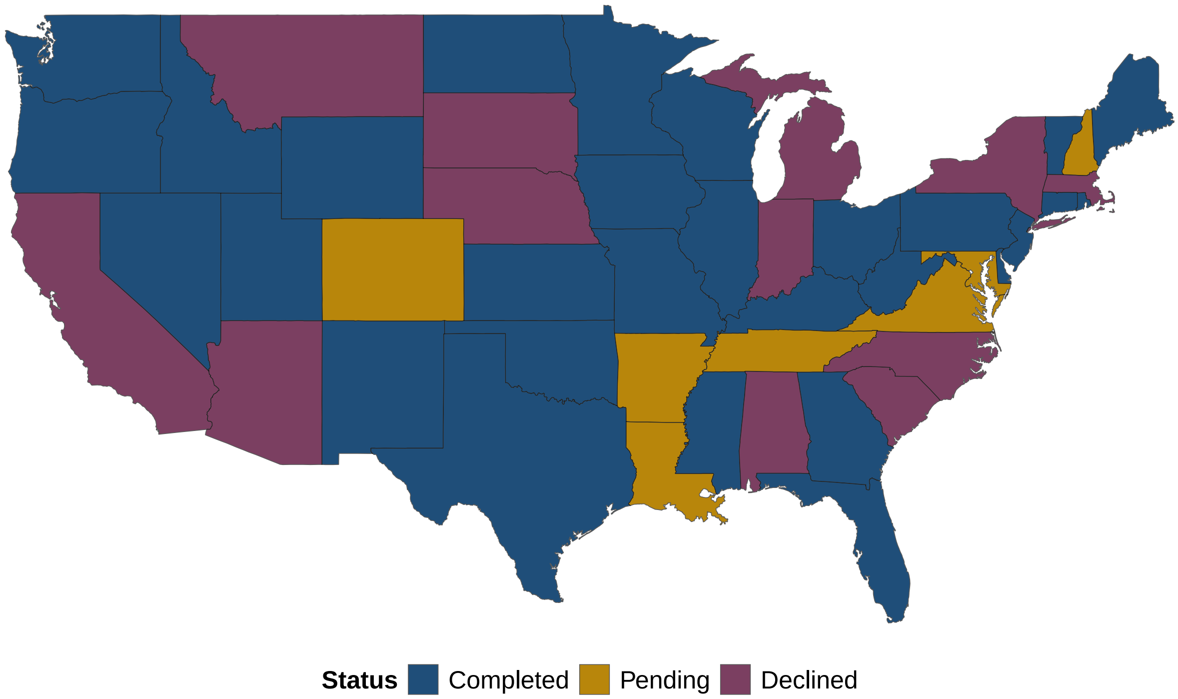

Figure 2. Status of Public Records Requests for Teacher Rosters: Green states provided complete official employee records for all districts (2017/18-2024/25). Orange states have pending data requests. Red states declined requests but data collection continues through other sources.

across states, we err on the side of false negatives and leave classifications as missing if position and assignment titles provide insufficient information.

Merge rosters to voter files: Next, we match all public school employees to the universe of all US registered voters between 2014 and 2025, provided by L2. As discussed in detail in Appendix C, we leverage probabilistic record linkage (Enamorado et al., 2019), and incorporate blocking on first name clusters, and where available, sex and year of birth. We then match on first name, last name, and date of birth, and retain all pairs with posterior match probability above 0.85. We then apply a geographic filter, where a matched pair is kept only if the voter's county lies within 150 miles of a district in which the employee worked. Finally, we enforce that each employee is matched to at most one voter (retaining highest-probability matches), and employees tied at the maximum posterior are dropped as unresolved. This prioritizes precision over recall and guards against false positive matches, particularly for employees with common names.

Identify employment changes and build panel: For each school year, we then code two types of turnover for each employee-year, recording whether an employee (1) moved between school districts or (2) left the state public school workforce completely, both relative to the previous year.[18] Our main analyses focus on district-switching and secondary analyses with exit are presented in Appendix E.

Descriptive Trends

We first present descriptive trends for both sides of our data pipeline, including the incidence and composition of culturally conflictual policies in schools over time, and the partisan and role-based composition of the teacher workforce across our sample.

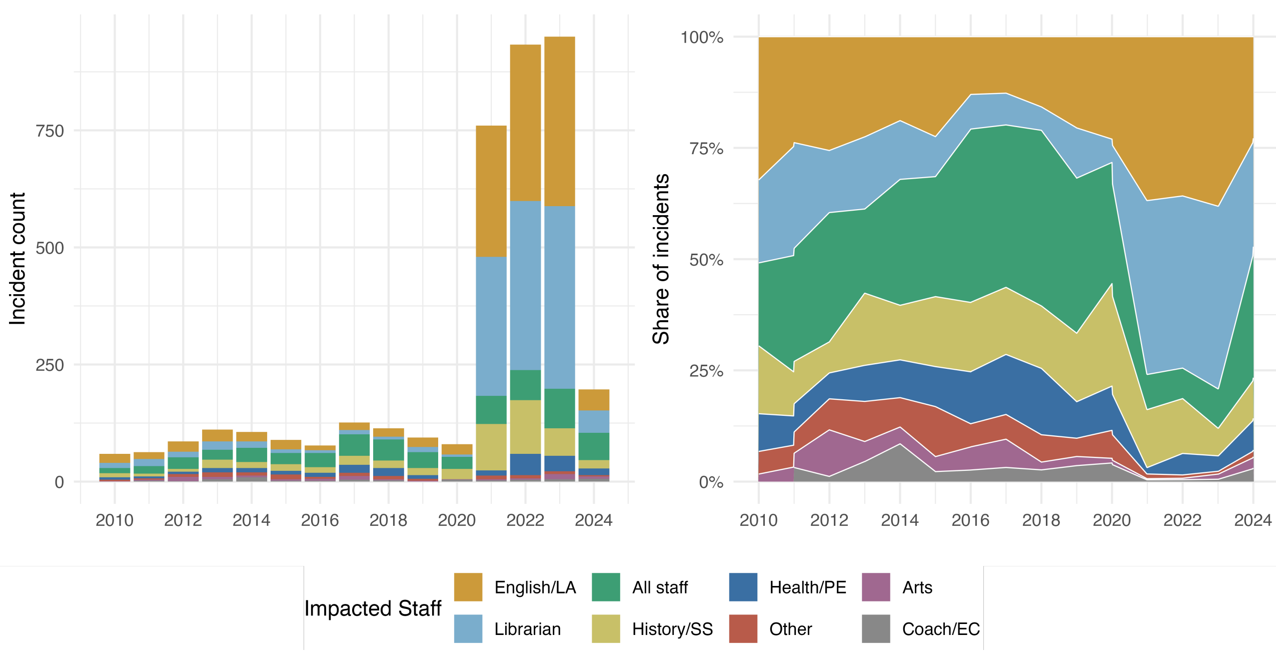

Figure 3. Both plots visualize district-level cultural policies from our three data sources. The left subplot shows the total count of policies over time. The right subplot shows the share of overall policies by which staff roles are primarily implicated, as shown by the fill colors.

[18]This second outcome is a bundle of teachers leaving the profession, moving to private schools or moving to a different state to teach. For some states, including Iowa, our administrative data allows us to decompose these outcomes.

Figure 3 visualizes how political conflicts in U.S. school districts have evolved over time and across staff roles. The left panel plots the total count of district-level cultural policies by year, pooled across our different data sources after de-duplication. Cultural policies have grown markedly since 2021, reflecting the wave of educational gag order legislation and book ban activity at state and local levels. The right panel shows the compositional shift in which roles bear direct exposure: policies that implicate librarians, English and language arts, and history and social science instruction has grown as a share of all cultural policies.

In terms of the workforce, employee partisanship varies substantially across the states in our sample, with the largest share of registered Democratic teachers in Pennsylvania and smallest in Idaho (see Figure 4a). We focus especially, however, on variation within states in our sample. Figure 4b illustrates average Democratic share at the district level in Georgia, revealing meaningful partisan sorting among public school employees. To quantify the amount of partisan clustering in educators' labor market, we compute Moran's I, a standard measure of spatial autocorrelation that ranges from -1 (perfect dispersion) to +1 (perfecting cluster), with 0 indicating an as-good-as-random spatial distribution. To account for sorting among the general electorate as the driver of employee sorting, we residualize the Democratic share in each district with the Democratic share among the general voter population.[19] Figure A4 visualizes this: each district is plotted against the mean residualized Democratic share of its contiguous neighbors. Partisanship of school employees is strongly spatially clustered (a raw Moran's I of 0.79 (p=0.001)): the positive slope suggests that districts with high Democratic teacher shares are surrounded by similarly Democratic districts. Conditioning on the local electorate absorbs about 21% of that, but a large, significant employee-specific clustering remains on top of the electorate's own geography (Moran's I of 0.586 (p=0.001)). Even net of both the electorate and state fixed effects, the sorting remains substantial (0.32 (p=0.001)), suggesting that spatial clustering is not reducible to within-state partisan sorting.

[19]Here, we simply use the last-observed partisan registration for each voter.

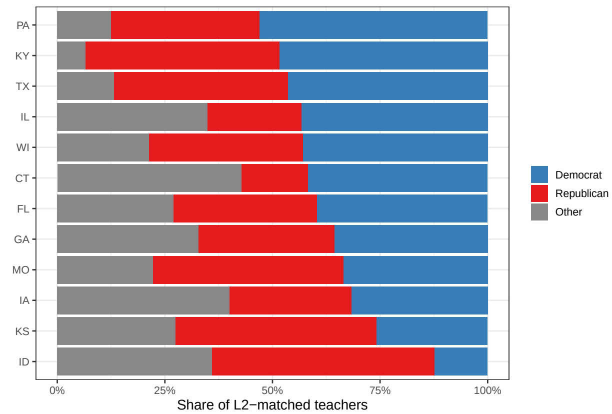

(a) Average school employee partisanship across states

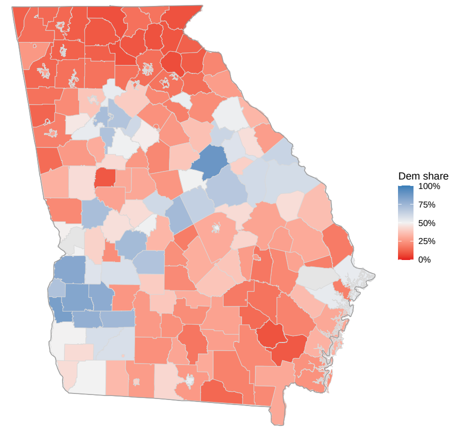

(b) Average employee Democratic share, Georgia

Figure 4. Teacher partisan composition varies substantially across and within states. Averages across districts with at least 20 L2-matched public school employees.

Finally, Figure 5 breaks down partisan composition by employees' role and subject areas. Among subject areas, ESL/bilingual, arts, English/language arts, and history and social science are the most Democratic while health and physical education are the least. Among educational roles, teachers at large have a Democratic share of about 40 percent while librarians are close to 35 percent. This figure shows descriptive insights on the joint distribution of exposure and alignment in our setting. The subject areas most directly targeted by the cultural policies we study are also among the more Democratic in the teaching workforce, suggesting that the workers who especially face restrictions on their professional discretion may also be more likely to be opposed to the direction of these policies.

Empirical Strategy

Estimands of Interest

Drawing on Table 1, our estimands of interest are the effect of enacting culturally conflictual policies (D_i) on an employee i's turnover among subgroups G_i defined by ideological misalignment

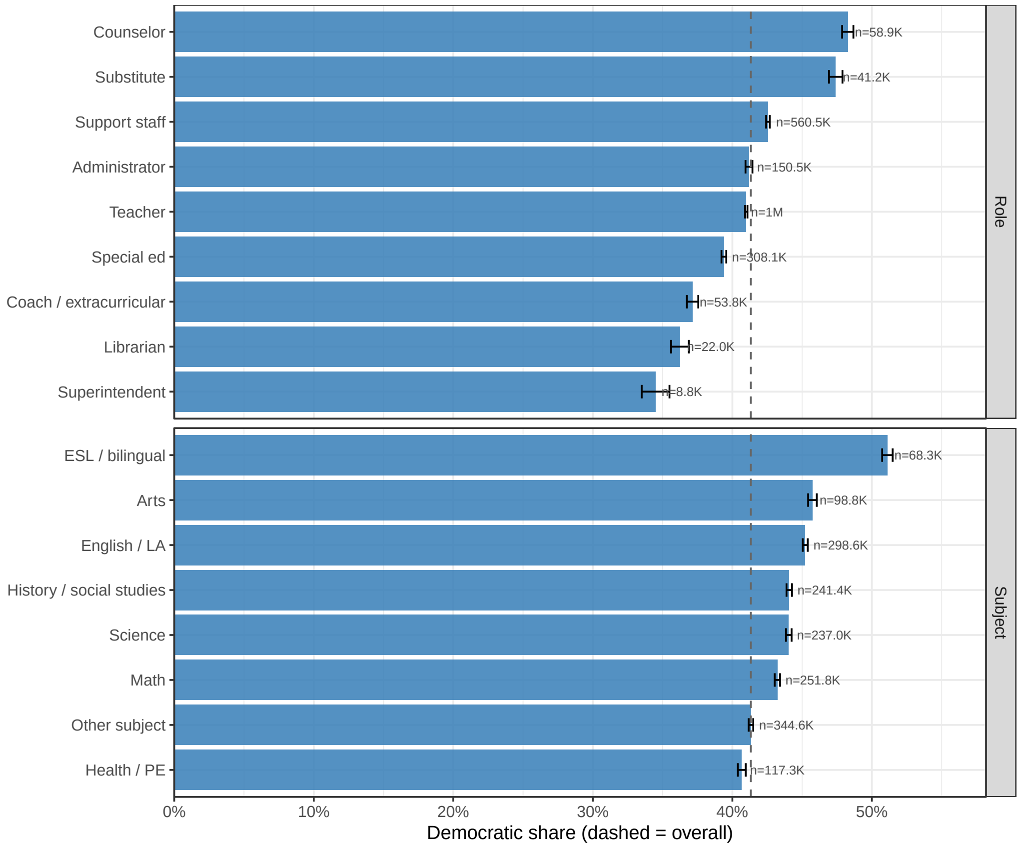

Figure 5. School employees' partisan composition varies substantially across roles and subject areas. Democratic share (among registered Democrats, Republicans, and Independents) among L2-matched educators by staff role and subject area. Sample sizes are shown for each category. The dashed vertical line marks the overall Democratic share across all matched teachers.

and exposure:

$$\tau_{ATT,g} = \mathbb{E}[Y_{it}(1) - Y_{it}(0)|D_i = 1, G_i = g] \tag{3}$$

For example, for G_i = misaligned, this quantity captures the average effect of exposure to cultural policies on an employee's turnover among ideologically misaligned employees. For this estimand to be identified for each group g, we require the usual assumptions of no anticipation and parallel trends to hold by group. For a given source of cultural policy (CRT restrictions, book bans, or CATO culture war incidents), we focus on the *first* instance in a given school district as the relevant treatment D_i.

To assess whether the treatment effects are larger for some groups than others—particularly ideologically misaligned as opposed to aligned employees—we further estimate the difference in the group-wise ATTs:

$$\delta_{G \in \{g_1, g_2\}} = \tau_{ATT, g_1} - \tau_{ATT, g_2} \tag{4}$$

This quantity then tells us how the average effect of exposure to cultural policies on turnover differs between individuals with different attributes (e.g., alignment). Importantly, we target an *effect modification* rather than a causal moderation, since we conceptualize both ideological alignment and subject-specific exposure as relatively fixed individual traits which cannot be meaningfully manipulated. This formalization means that, in our empirical design below, we do not need to assume mean independence of the potential outcome changes, $\Delta Y_i(g, d)$, and the group assignment G (Xu et al., 2026), to estimate these estimands. That is, we acknowledge the fact that characteristics such as partisanship and job exposure are endogenous bundles, and we are not asking whether being a Democrat *caused* a different response to the policy. Instead, we target the more policy-relevant descriptive quantity: the difference in the effect sizes across different types of employees.

Empirical Design: Employee-to-Employee Matching

To identify the causal effect of cultural policies, one could use a staggered DiD design, comparing employees in treated to those in untreated school districts. However, since these policies are not randomly assigned to districts, employees in treated districts may differ from those in control districts in ways that could determine their tenure and willingness to remain in a given district following the policies, e.g., partisan composition, average experience, or student enrollment. Instead, to make the key identifying assumption more plausible, we use a matching procedure that leverages our administrative data of 2.9 million employees to match each treated employee with an employee in a *never-treated* district in the same state who is comparable across a large set of observable covariates. Specifically, for each of our 140,000-160,000 treated employees—depending on the source of conflict—we find a pair employee in a control district who matches exactly on the following administrative covariates *at the time of treatment onset*: partisan alignment with the policy, our 18-category position classification described above, grade-level assignment, full-time employment status (e.g. full-time vs. part-time), binned years of in-district work experience (e.g. 20-25 years in this district), state-specific binned age (e.g. 24-29 years old), and state (e.g. Oklahoma). Further, we also pair employees from similar *school districts* by matching on district-level covariates with Mahalanobis distance: district-level student enrollment (in the year of treatment) and per-pupil spending (in 2019-20). To ensure the best possible match for each treated employee, we match with replacement. This implies that our key group-indicators (ideological alignment and exposure) are fixed within pairs and any within-pair DiD estimates target the group-specific ATTs. We describe our matching procedure in much more detail in Appendix D, where we also report diagnostics to show that our matching successfully reduces observable covariate imbalance.

The primary benefit of our matching approach is that it allows to estimate differences between very similar pairs of employees who primarily differ in their treatment status – for example, a change in turnover propensity between teachers with the same position, experience, employment status, age, working in very similar school districts in the same state, yet differ on whether their employer adopted an anti-CRT policy or not. Additionally, by calculating treatment effects *within matched pairs*, we solve the issue of negative weights and effect heterogeneity that often plague staggered DiD designs.

Table 2 reports the matching yield across our treatment sources. The treated pool is every employee who was under a cultural policy and who satisfies the two-year pre-treatment history requirement, and the yield is the share of them for whom an exact-cell control exists. With about 90%, yield is high but not complete. Controls are used an average of 3.4 to 4.9 times due to replacement, but about 50% of controls are uniquely matched to a single treated employee.

Table 2. Matched pairs, by treatment source, twelve states. *Controls* is the number of distinct control employees used; *Reuse* is matched pairs per distinct control; *Used once* is the share of controls that serve exactly one treated employee. Matching is 1:1 with replacement, so a control may be paired to several treated units.

Source	Treated	Districts	Pairs	Yield	Controls	Reuse	Used once
CRT ban	139,695	112	127,026	90.9%	25,680	4.9	50.8%
Book ban	158,903	168	138,598	87.2%	30,946	4.5	52.7%
CATO (direct)	160,089	283	143,812	89.8%	42,333	3.4	55.4%

Estimation Strategy

Once we have our matched employee-year panel data set, we estimate the following regressions using OLS:

$$Y_{ipt} = \beta_0\, Post_t + \beta_1 \left(D_i \times Post_t\right) + \alpha_{p(i)} \;+\; \gamma_t \;+\; \delta_i \;+\; \varepsilon_{ipt} \tag{5}$$

$$Y_{ipt} = \beta_0\, Post_t + \beta_1 \left(D_i \times Post_t\right) + \beta_2 \left(Post_t \times G_i\right) + \beta_3 \left(D_i \times Post_t \times G_i\right) + \alpha_{p(i)} + \gamma_t + \delta_i + \varepsilon_{ipt}, \tag{6}$$

where Y_{ipt} is the turnover across districts or out of the profession for employee i in pair p in year t, $Post_t$ indicates years following the enactment of the policy, D_i indicates whether a pair member was treated by the policy ($D_i = 1$) or in a never-treated district at the time of the treatment ($D_i = 0$), and G_i are the group indicators for ideological alignment and/or exposure. $\alpha_{p(i)}$ are matched-pair fixed effects (to isolate comparisons between paired teachers), γ_t are year fixed effects (to account for unobservable patterns over time, such as specific post-COVID turnover trends), and δ_i are employee fixed effects (to account for unobserved employee-specific characteristics that determine turnover behavior). The group-specific ATTs from (3) are then estimated with β_1 (for $G_i = 0$) and $\beta_1 + \beta_3$ (for $G_i \neq 0$). ε_{ipt} are clustered by the school district at treatment assignment. For the treatment group, we restrict our analysis to employees who already worked in the district *when the policy arrived*, i.e., we drop 4.9% of employees who moved into an already-treated district. These movers pose both mechanical problems (their move inflates pre-treatment turnover estimates automatically) and selection issues (the policy was already public when they chose that district, so treatment is self-selected, not assigned).

Results

Table 3 shows the DiD and triple DiD estimates across the different policies, using an employee's move across school districts within a state as the main outcome. Similarly, Figure 6 depicts the marginal group-specific DiD estimates for aligned and misaligned employees, both pooled and across exposed and non-exposed public-school staff. As expected by our theory, we do not find increased average turnover among *all* staff as a result of the cultural policies (columns 1, 3 and 5). Yet, we do find a bifurcating effect on turnover when distinguishing misaligned and aligned employees. For CRT, aligned employees in treated districts were about 2 percentage points (p=0.098) *less* likely to leave their district compared to educators in control districts post-policy. Misaligned employees, in contrast, were 1.7 percentage points (p=0.048) *more* likely to leave their district after the policy took effect relative to matched control employees, leading to a 3.7 percentage point difference in ATTs between aligned and misaligned employees. Given a baseline turnover rate between districts of only 3% pre-treatment, these effects are politically substantial. Additionally, and again in line with our theory, we find that this bifurcating effect is strongest among exposed bureaucrats (with a triple DiD estimate of 0.04 among exposed compared to only 0.009 among non-exposed staff, left panel in Figure 6).

For the other two treatment measures, the estimated effects are directionally in line with our theory, but substantially smaller and more noisy. For example, for Cato, we again find a significant difference in turnover propensity following the policies among exposed employees between misaligned and aligned bureaucrats (triple DiD of 0.016, p=0.01). Particularly, misaligned and exposed employees were 1.1 percentage points (p=0.03) more likely to leave their district, while aligned and exposed bureaucrats were slightly, although insignificantly more likely to stay following the policies (-0.5 percentage points, p=0.49, Figure 6).

For book bans, however, there is little evidence that they drove bureaucrats to reallocate across districts. Instead, when disaggregating the book-ban exposed employees into their subgroups, we see that these bans pushed misaligned, but not aligned *librarians* to completely exit their profession (triple DiD of 0.10, p=0.01, Figure A6). Misaligned librarians exited at a 5 percentage points higher rate following the treatments, while aligned librarians were 5.1 percentage points *less* likely to leave. With a turnover rate of 8% among control district librarians, these estimates are again meaningful. This is consistent with book bans being more targeted at libraries and spatially clustered than CRT restrictions, often driven by state-wide legislative changes: neighboring districts are markedly more likely to share an enacted book ban than an enacted CRT restriction (join-count z-score, which estimates spatial autocorrelation with binary data, is 0.26 vs. 0.18, both $p < 0.001$). Moving to a neighboring district within a state is therefore unlikely to be sufficient to escape book bans, so leaving the profession may be the only option for misaligned employees.

Ultimately, to assess the plausibility of the parallel trends assumption and the timing of the treatment effect, Figure 7 shows the marginal DiD estimates for aligned and misaligned staff, relative to the onset of CRT restrictions. Reassuringly, pre-treatment trends closely track one another for treated and control school employees in both groups, but start to diverge in opposite directions soon after treatment onset.

Table 3. Effect of enacted culturally conflictual policies on district move

Dependent Variable:	Outcome					
Treatment	CRT ban		Book ban		CATO	
	DiD	Triple	DiD	Triple	DiD	Triple
	(1)	(2)	(3)	(4)	(5)	(6)
Variables						
post	0.0027	0.0193**	0.0046	0.0055*	0.0017	0.0030
	(0.0031)	(0.0088)	(0.0031)	(0.0031)	(0.0030)	(0.0044)
treat × post	-0.0023	-0.0207*	-0.0024	-0.0009	0.0024	0.0016
	(0.0065)	(0.0126)	(0.0033)	(0.0038)	(0.0033)	(0.0051)
post × misaligned		-0.0342***		-0.0039		-0.0034
		(0.0127)		(0.0057)		(0.0040)
treat × post × misaligned		0.0374***		5.31×10^{-5}		0.0060
		(0.0132)		(0.0061)		(0.0049)
Fixed-effects						
pair_id	✓	✓	✓	✓	✓	✓
year	✓	✓	✓	✓	✓	✓
teacher_id	✓	✓	✓	✓	✓	✓
Observations	1,406,876	971,598	1,391,115	898,413	1,254,766	765,926
R^2	0.22472	0.22149	0.25712	0.25634	0.25181	0.25070
Pre-treat. control mean	0.0304	0.0323	0.0203	0.0205	0.0255	0.0262
Num. pairs	120,806	82,705	132,711	85,376	120,648	73,988

Clustered standard-errors at the district-at-treatment-onset in parentheses
*Signif. Codes: ***: 0.01, **: 0.05, *: 0.1*

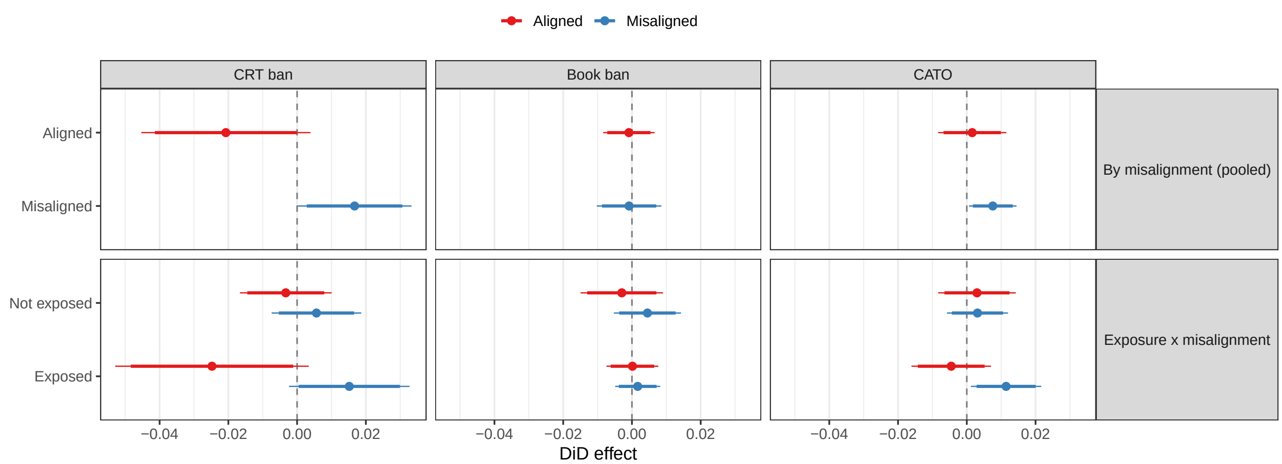

Figure 6. Marginal DiD estimates on district move, across subgroups and treatment measures, with 90% and 95% confidence intervals

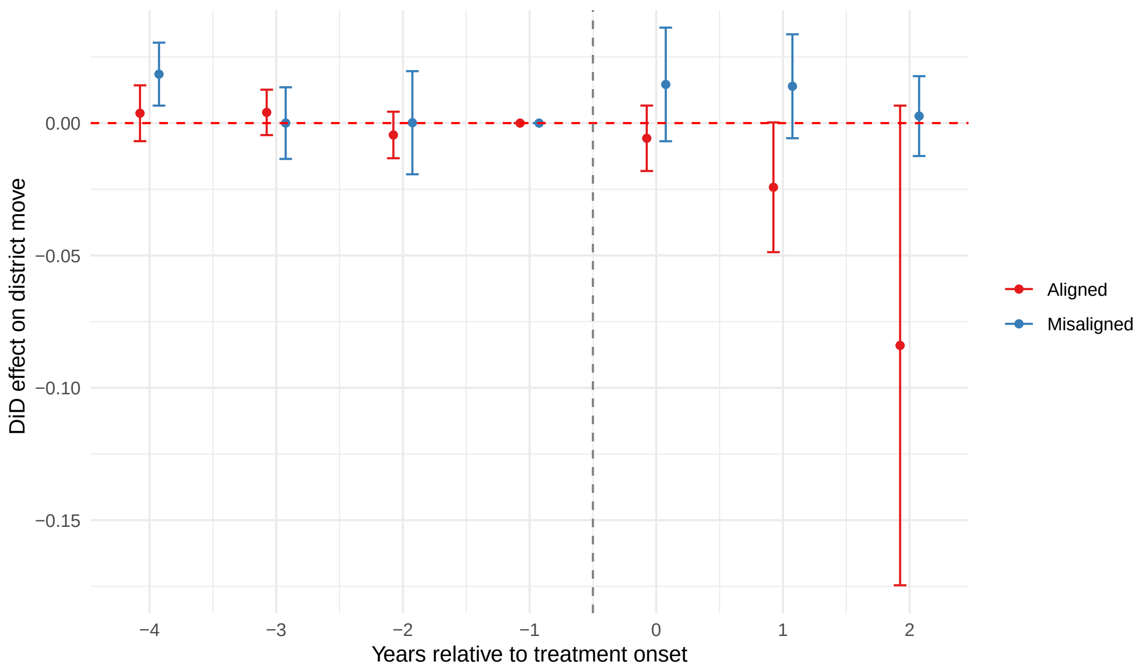

Figure 7. Event study graph for the marginal DiD effect on district move (CRT), with 95% confidence intervals.

Additional Findings and Robustness Checks

Characteristics of Staff Who Are Switching or Leaving

To further contextualize how the gains and losses in school staff is distributed across types of employees and districts, we split the DiD for aligned and misaligned individuals by observables. Figure 8 shows the treatment effects by age terciles. Interestingly, the youngest tercile shows the widest gap, with treatment effects of about 3 percentage points for misaligned vs. about -4.7 percentage points for aligned staff. Both treatment effects shrink toward zero for older cohorts. This suggests that the sorting effect of cultural policies is concentrated among early-career educators, who likely have the least discretion to begin with and higher mobility. This also speaks to our theory's implication that discretion and alignment are substitutes: because the value of alignment is greater for bureaucrats with less discretion, the gap between misaligned and aligned teachers should be widest among those with least baseline latitude.

Conversely, Figure 9 shows that poorer school districts (measured as the share of students receiving free or subsidized lunch) bear most of the additional turnover among misaligned employees, while more affluent districts do not lose misaligned staff and are better able to retain aligned workers. Figure A9 shows a similar concentration of the sorting effect among districts with larger shares of non-White students.

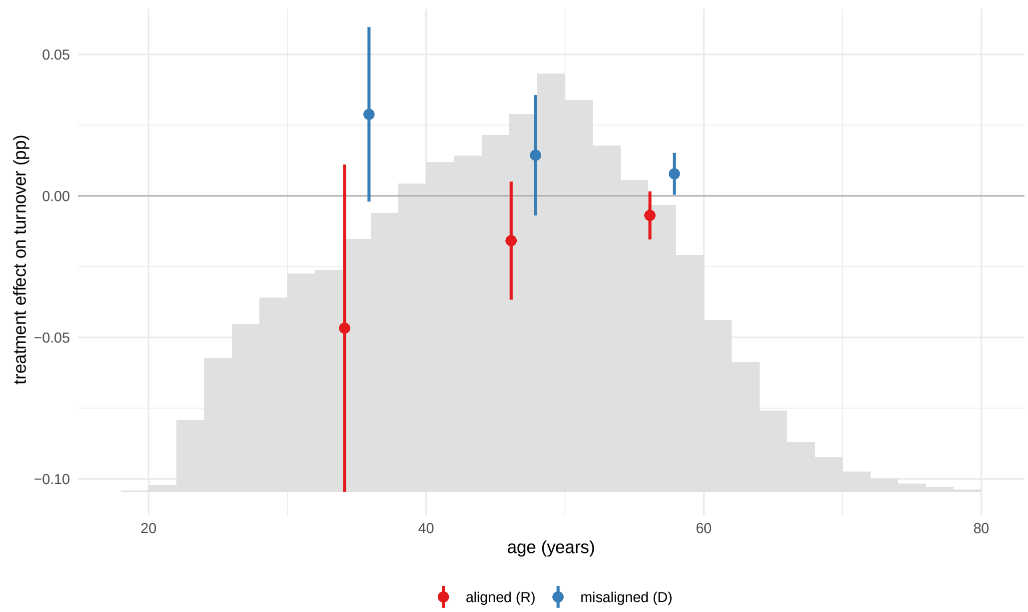

Figure 8. Alignment-specific treatment effect by age tercile on district move. Moderator distribution shown in grey.

Where Are Treated Staff Moving To?

Additionally, we analyze where treated employees are relocating to. As Figure 10 shows, treated movers overwhelmingly land in non-conflict-affected districts (80%-89%) after onset, and at higher rates than before. These post-onset moves are also less financially rewarding. Moves made before onset raise average salaries by 7-10% on average, while transitions made after conflicts yield only 0-5%, for both misaligned and aligned staff alike (Figures A10 and A11). If misaligned workers were simply chasing better offers, their post-move salaries should diverge from those of aligned staff, or their destination choices should track pay. Since we observe

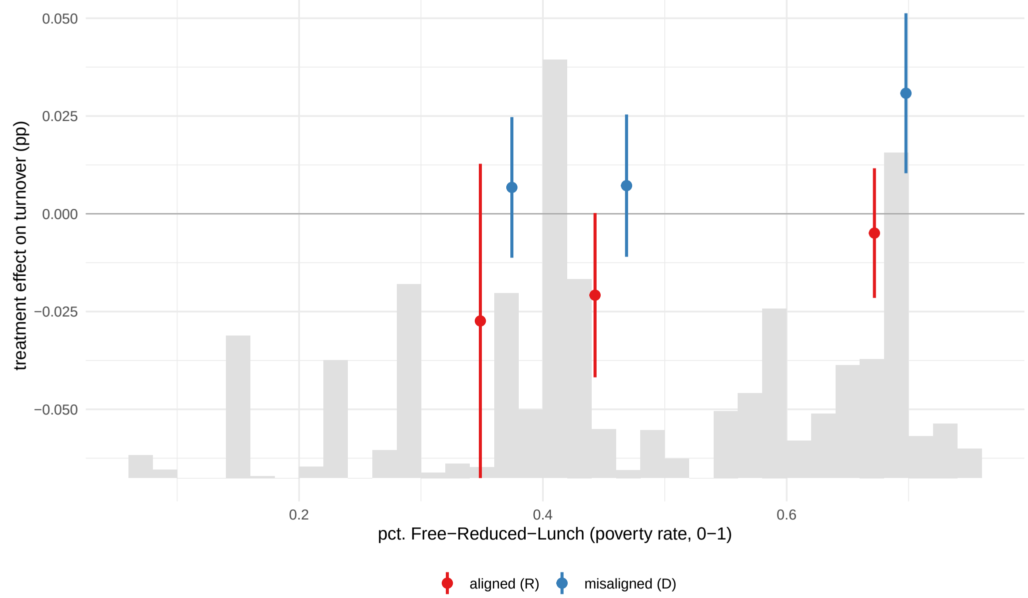

Figure 9. Alignment-specific treatment effect by poverty tercile on district move. Moderator distribution shown in grey.

neither, their relocation seems to primarily be a response to conflict and does not allow them to substitute material and non-material benefits.

Robustness

Our estimates are robust to various sensitivity analyses. First, since our results heavily hinge on the measurement of alignment, one maybe be concerned that biases in the ability to match education bureaucrats to voter files could affect our inferences. To assess how much linkages with L2 diverge from a random ideal, we compare demographics recorded in our administrative data between employees who were linked and those who were not (including race, sex, salary, experience, degree, etc.). Linked employees are indeed somewhat more White, more female, and more senior across different roster states. However, as we show in Figure A1, only 19% of all covariate comparisons exceed the conventional threshold of a standardized mean difference of 0.1, and none is above 0.5. What is more, when we reweight our matched sample using inverse-probability weighting to better represent the full roster in terms of observables, we find

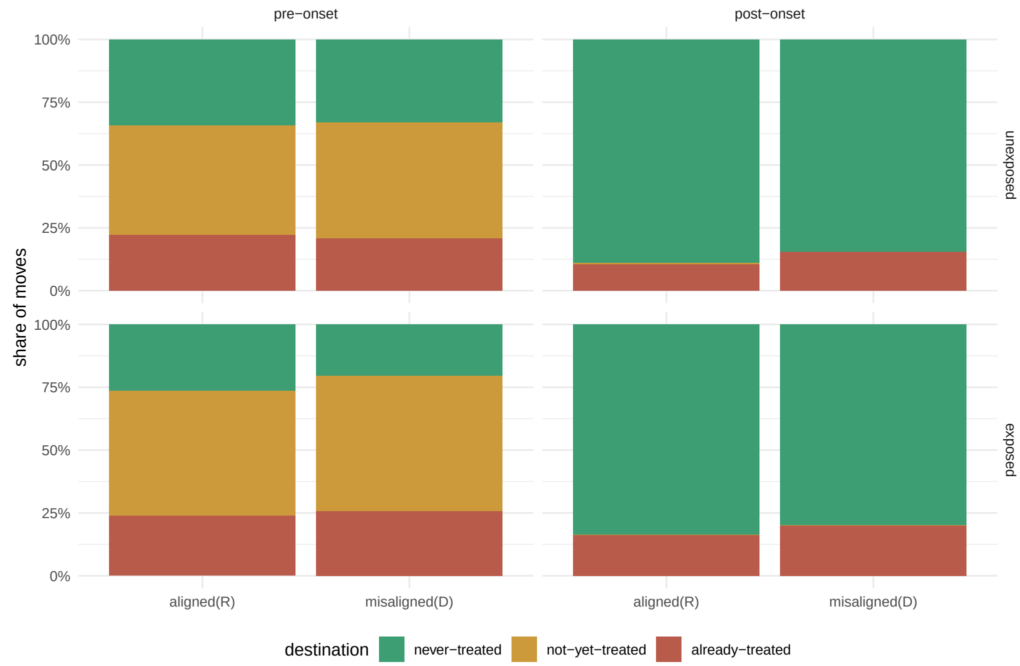

Figure 10. Destination by exposure and alignment of district movers (by destination CRT status). Controls omitted by construction.

almost identical point estimates for the triple DiD among aligned and non-aligned bureaucrats (Table A8).

Second, one may be concerned that the differential trends for aligned and misaligned bureaucrats in our main specification could simply reflect a secular trend, within certain districts, certain subjects or certain partisan groups. However, as Table A10 shows, the estimates for district move are almost identical when comparing aligned and misaligned employees in the same district (column 2), in the same subjects (column 3) or with the same post-treatment partisan registration (column 4).

Third, we rule out that the remaining imbalance in district size after matching might confound our results. Again, our estimates are robust to adding controls for student enrollment in addition to our fixed effects (Table A11). Hence, our estimates are unlikely to be an artifact of differential trends by district size.

Fourth, we assess whether our results are driven by administrative staff outside the classroom rather than teaching staff, that are more core to our theoretical argument. Reassuringly, the estimates are very similar, or even larger, among teachers only (Figure A12).

Finally, we look at whether misaligned employees are also disproportionally more likely to switch *schools in the same district* in response to cultural policies. We see this as a placebo outcome, since transferring between schools does not help to escape a district-wide policy. Indeed, we find no evidence that misaligned staff is more likely to leave their specific school following the treatment (Figure A13). If anything, they are less likely to do so (for Cato), which is likely a result of district and school moves being mutually exclusive.

Qualitative Evidence

Interviews with Public School Employees

We complement our quantitative analyses with qualitative interviews with 10 current and former K-12 employees. This section pursues three objectives. First, we examine whether educators' own accounts are consistent with the two channels of alignment and exposure our framework specifies. Second, we use interview evidence to provide concrete detail of how our framework and findings manifest in real-world practice. Third, we use interviews to surface mitigating forces that may plausibly explain gaps between the costs we theorize and the turnover we observe, including why some employees depart from what our main specification predicts. The first two goals are served by organizing evidence around the exposure-by-alignment cell a teacher falls into while the third is better served by case type (typical vs. deviant, explained below) and the specific constraints a teacher describes, since these cut across cells.

Our district selection from which we interview public school employees follows a mixed-methods residual-based case selection strategy (Seawright, 2016; Nalewajko, 2025; Dunning, 2012): we compute district-level residuals from the main triple difference-in-differences specification,[20]

[20]We use the anti-CRT policies as the main treatment. Further, district residuals can help to capture

classify districts as "typical" or "deviant" based on their minimal or maximal distance from model predictions (i.e. rates of employee turnover that are much higher or lower than our model predicts), and cross this classification with treatment status. Typical cases are useful for process tracing mechanisms, allowing us to observe whether the theorized channels are in fact operating in the cases where the model fits well. By contrast, deviant cases identify districts where turnover substantially diverged from predicted values, pointing to conditions that may not be captured by the main specification and revealing potential scope conditions. We select 19 districts across Florida, Georgia, Kentucky, and Texas, with additional criteria ensuring variation in district size (based on enrollment), conflict intensity, and contact feasibility.[21]

Mirroring a process shared by other applications of this method (Nalewajko, 2025), this classification operates at the district level, and is worth distinguishing from what it implies about any individual respondent. A typical-treatment district is one where the model's aggregate prediction fits well, not one where every individual educator's outcome matches what her own cell predicts. Because the treatment is a district-level policy and case selection operates at the district level, residuals are constructed by aggregating the model's individual-level predictions to the district before differencing against observed turnover. For example, an exposed, misaligned respondent who remains in a typical treatment district is, at the individual level, a case our theoretical framework would flag as elevated-risk; her retention is not evidence against the theory so much as the kind of gap the evidence on mitigating forces below might explain.

Within each district, we identify public-school employees from the state roster, prioritizing those in the matched analytical panel and stratifying by role-based exposure and ideological alignment. Within each district, the four exposure-by-alignment cells carry different evidentiary weight. In following our theoretical two-by-two, cells A and B (exposed, misaligned/aligned) help to capture divergence between rising costs among misaligned employees and improved retention

district-level variation in outcomes that the model cannot attribute to matched covariates.

[21] 6 districts are 'typical' treatment, 6 are 'deviant' treatment, 4 are 'deviant' control, and 3 are 'typical' control. The interviews so far cover 5 of the 19 districts as well as an interview with an educator in Virginia through snowball sampling.

among aligned employees. Cells C and D (non-exposed, misaligned/aligned) can isolate the alignment channel alone, holding discretion fixed. Full details on district selection, sampling strategy, outreach protocol, and interview questions are in Appendix F.

Evidence for the theorized channels

First, we find that public school employees consistently report recognizing contemporary policy conflicts like CRT as interventions that reduce their own discretion. These preliminary interviews are suggestively consistent with our framework. Long-term teachers across several interviews independently framed the current period as distinct from ordinary political cycling. A teacher with nearly three decades in the classroom directly said: *"historically, trends come and go. You know, it gets a little more conservative, it gets a little more liberal, but this trend is not looking like it's going away...they're digging deeper and deeper into state control of the schools and the policies."* A science teacher in her fourth decade of teaching made a similar comparison: *"I have seen our governor take more control of schools than I've ever seen a governor take control."*

The alignment channel — costs that follow from the relationship between a employee's personal ideology and the institution's official position, independent of whether a teacher's own discretion is regulated — appears across interviews regardless of subject area and in both directions. One misaligned and exposed educator described her own earlier departure from a prior position as a deliberate choice driven by how *"'the person in charge of library services was never gonna support diverse books."* Another explicitly non-exposed but misaligned teacher discussed how *"it's swung so far that now you can't mention anything in [state], anything culturally diverse or slavery, or any of it. ..we in the public school system...raise children with the intention of being a good citizen, a tolerant citizen. All of these things that we as a society have deemed important have gone out the window. And now, those are bad things...if you have those thoughts, you are woke in a bad way."* That the alignment cost registers even where a teacher's own domain is not the policy's direct target is consistent with our claim that the cultural policy shifts a for the entire workforce.

We also have evidence in the opposing direction. An aligned interviewee in an administrative role in a treated district reported unusually low turnover and a surplus of job applicants for open positions, attributing this in part to the state's and district's policy direction and their compliance with new requirements: *"one of the ways to measure morale of the staff...is teacher retention...we now have more applicants than we need for every position...that's a positive reflection on the way things are trending here."* The educator located this in cultural alignment between the school and community: *"we're in a conservative community...we're trying to reflect the community, we're not pushing the envelope"* and noted that *"the cultural fits are very good"* among new hires. The same respondent reported no known departures due to cultural policies. A second aligned respondent, teaching a non-exposed subject in a control district was explicit about her own alignment with the policies' direction — *"critical race theory is something I'm 100% opposed to"* — yet reported that the broader political climate *"doesn't really touch my classroom."*[22]

Directly exposed, misaligned teachers layer a discretion cost on top of the alignment cost. One teacher first situated this generally, stating *"it's interesting that politicians decide how we'll teach and what we'll teach, and they aren't teachers...they shouldn't be deciding those things..[or] what I do,"* before giving a concrete instance from her own tested standard: *"I teach science. I teach weather. I teach climate. And I can't say climate change....there was a very specific question about climate change on the [state] science test for 8th grade...I wasn't allowed to teach it. And yet, it is still a [state] standard."* Another described the same mechanism operating directly on her own classroom judgment: *"I want to call students whatever they want to be called, and we're not allowed to do that, and that makes it difficult."* A third teacher, whose classroom library

[22]We also asked misaligned interviewees directly whether they perceived colleagues as supportive of the enacted policies. These answers are informative but secondhand, and should be read as characterizing perceived alignment. One teacher noted that *"I would say the very, very conservative teachers have [been supportive], and every administrator has because...they know how to yes ma'am, yes sir...they'll do whatever they're told to do."* A second teacher discussed perceived support of curricular retrenchment: *"I know people who...like current administrations, both district, state, and national...typically they tend to support that with, we need a more conservative way of ...[getting] back to the basics."*

falls under book-review requirements, described the discretion cost in practical terms: *"every book we keep in our room has to be approved by somebody somewhere"* – a burden that is severe enough that among her co-teachers *"the [classroom] bookshelves are empty."* The same teacher described the enforcement machinery our framework treats as compounding the discretion cost beyond the restriction itself, reporting self-censorship around colleagues of unknown political orientation: *"you have to be very aware of who you're talking to...[when we know] who is in the room, we know not to say the trigger words."* Several teachers further noted a form of psychological perceived reduction in discretion and respect for their expertise. One teacher said *"our governor decided that anyone that's been a veteran, or their spouse...without a degree, can be a teacher...What he's saying is that you guys don't matter, we could hire anyone...to do it. And that's very disrespectful,"* while another claimed *"it seems like people don't see teaching as a profession anymore and [see it more like] a service...it's very disheartening. I would ...not recommend people to go into teaching."*[23] By contrast, the aligned and non-exposed respondents did not report any such costs.

Mitigating forces and the cost-turnover gap

Interviews also speak to the gap between underlying cost and observed turnover. While many teachers relayed that cultural policies imposed high costs on the symbolic benefits they gain from teaching—consistent with our theory—many also remained employed in their districts because of the difficulties of switching employers. These constraints become especially important as employees near retirement eligibility and pension vesting: *"I'm too old to leave right now...if there wasn't light at the end of my tunnel, I don't know [what I'd do]."* The same teacher described a colleague's decision in explicitly quantified trade-off terms: *"she has a husband who's still working, and decided that the amount of money she would lose was worth not staying for 2 more years of aggravation, so she's retired this year."*

[23]Teachers in public meetings have invoked similar concerns about a sense of disrespect toward the profession. They have also vocalized both alignment and misalignment with anti-CRT bans. See Appendix H for additional qualitative evidence from *DistrictView* (Holman et al., 2025) public meeting records.

Financial and housing constraints also trap teachers who would otherwise leave or tie them to a district they might otherwise leave for political reasons. One teacher described a colleague who *"would prefer to pack up and go somewhere else, but really can't afford to…due to her family's situation...it's real hard to cold turkey say, I'm gonna quit and do something else, when you have to pay the mortgage."* Another described a younger colleague who disliked the local political climate but stayed in part because she lived rent-free in a comparatively high-paying district. Others describe colleagues sorting toward private schools that trade job security and pay for restored discretion — the w for v substitution in our utility framework: *"...I personally know 3 people who have gotten out of teaching public school because of what's going on …one of them left teaching altogether, the other two went to private school where they don't have to put up with that... but you know there are tradeoffs – they don't have benefits, they don't have a union, they can be fired any minute, you know, and…the charter and private schools pay less."*

By contrast, one misaligned exposed educator explicitly cited a *political* reason holding them back from quitting: *"one of the reasons people stay is because...to go back to some of the teachers who were teaching marginalized people's perspectives, or work by marginalized groups...they know that whoever comes in behind them is not going to do that. I mean...I think the plan is to drive us all out and replace us all with 22-year-old [ideologically-aligned] university grads."*

Not every account supports a cultural-policy-specific story. One respondent in a deviant treatment district described their departure due to an unrelated generic state personnel statute unconnected to curriculum content in one case. Another teacher described departure due to an interpersonal dispute with school leadership. Several respondents across cells and case types also raised student behavior and discipline management, unprompted, as a large source of dissatisfaction beyond any specific culturally conflictual policy. Additional evidence is in Appendix G. We report these not as evidence against our framework, but as a reminder that our design does not claim cultural policy is the only or even the dominant driver of any given departure—indeed, beyond the issues mentioned here, teachers also noted challenges from pandemic disruptions, device policies, and emergent AI in the classroom as challenges to morale. Our study evaluates the additional effect of cultural policies, net of these other forces.

Discussion and Conclusion

This paper develops and tests a theory of workforce sorting under cultural conflict. We argue that cultural policies impose uneven costs on institutional workforces not merely by constraining professional discretion, but by making official claims about whose professional values count as legitimate within the institution. Using administrative rosters covering 2.9 million public school employees across twelve states, linked to district-level policy adoption and to voter registration records, we find that enacted restrictions on the instruction of Critical Race Theory shifted the workforce spatially across school districts along ideological lines, and book bans split the workforce by differential retention of librarians.

Our central theoretical insight is that cultural policy does not purely impose a burden on a strained workforce. Instead, it makes a claim on value legitimacy that affects bureaucrats differentially. Existing accounts of how political shocks affect public workforces are largely unidimensional. Red tape, compliance costs, sanctioning risk all raise the cost of remaining and predict attrition (Bandiera et al., 2021; Honig, 2024). Similarly, empirical work on ideological alignment between bureaucrats and institutions primarily focuses on *misaligned* workers (Richardson, 2019; Bolton et al., 2020; Spenkuch et al., 2023; Brown et al., 2026). In contrast, we argue and show that the same enacted policy that raises departure rates among ideologically misaligned employees can lower them among their aligned colleagues, leaving average turnover essentially unchanged.

The sorting we demonstrate is consequential for several reasons. First, unlike bureaucrats' temporary behavioral adjustments to ideological misalignment on the job (Spenkuch et al., 2023; Wirsching, 2026a), employment changes are more extreme and durable decisions that reshape the public sector workforce.[24] This means a transient ideological school board majority can enact cultural policies that leave a lasting imprint on an institution's personnel long after the conflict is settled or repealed.

Second, the losses are not evenly distributed. Departures concentrate among early-career employees and in districts serving the poorest and most diverse student populations, which are already districts with the largest difficulties to retain teaching staff (Sorensen and Ladd, 2020; Ingersoll, 2001; Jackson, 2009). Moreover, such sorting can be consequential for the democratic and pluralistic functions of schools (Gutmann, 1999; Dewey, 1916; Nelsen, 2023; Tiebout, 1956). Teachers shape students' educational outcomes (Kraft et al., 2016; Chetty et al., 2014) and influence their political socialization and civic development (Sapiro, 2004; Jones, 1971; Geerlings et al., 2019). When educators move based on their willingness to (dis)engage with cultural content that often concerns minoritized groups, the consequences can compound those of the policies themselves. Restrictive curricula constrain what is taught, but misaligned teachers who remain can moderate these effects – for example through contextualizing restrictions, maintaining relationships with students whose identities are implicated, or simply signaling that dissent is possible. Sorting can remove these educators, producing ideologically insulated learning above and beyond what these policies alone generate, with potential downstream differential effects on students' tolerance, civic norms, sense of belonging, and political development (Haas and Lindstam, 2024; Bonilla et al., 2021; Nelsen, 2023).

Lastly, our findings speak to a broader point about the politicization of public institutions. Some public employees benefit from civil service or tenure protections that insulate them from politically motivated hiring and dismissal. But such protections govern what an employer may do to a worker, not whether the conditions of the job become costly enough that misaligned workers exit on their own. As a result, even for policies that are not about personnel at all, politicization can arise *endogenously*: it emerges from the aggregation of individually voluntary decisions

[24] The average employee in our sample works in their district for 13 years (median 10).

rather than from any deliberate program of political screening or employment discrimination, and it is therefore largely invisible to the legal and administrative machinery built to prevent it. For example, much attention has been paid to the Trump administration's most visible, contestable personnel interventions (e.g., mass layoffs of agency staff or reclassification of career positions) (Jaffe et al., 2025; Shao and Wu, 2025). Far less visible is the sorting induced by policies that never touch employment at all, which leave no dismissal to contest and no aggregate personnel change to detect it by, yet can lead to harmful downstream consequences, including widening affective polarization and chilling public discourse (Altman, 2025).

Although our case and findings focus on public school employees, the mechanisms we identify likely operate across professional institutional settings where cultural policies make official claims about whose professional judgment is legitimate. For example, news accounts document medical residents avoiding placements in states with abortion restrictions, with OB-GYN programs facing particular pressure as the restrictions most directly implicate their clinical domain (Rovner and Pradhan, 2024). In the legal profession, associates at BigLaw firms that struck deals with the Trump administration resigned in growing numbers, citing the mismatch between their firms' institutional positions and their own professional commitments (Weiss, 2025). Whether the role-based exposure and ideological misalignment logic we identify here travels to law, medicine, or other employees like university faculty or federal agency staff navigating culturally-conflictual mission changes or shifts in institutional neutrality is a fruitful question for future research.

While this paper focuses on employment change as the observable consequence of institutional misalignment, cultural policies likely produce a range of behavioral responses among employees who remain. For example, the conditions under which workers exit versus exercise voice, resist, or stay and adapt (Hirschman, 1972) is an important priority for future research. Recent reporting on university faculty subject to syllabus surveillance laws and organized pressure campaigns from groups like Turning Point USA documents strategic adaptation rather than departure: professors maintaining two versions of course materials, avoiding keywords likely to attract scrutiny, or narrowing content to avoid formal complaint (Patel, 2026). Additionally, while our study focuses on current employees, future research could also examine whether culturally conflictual policies shape who enters a profession in the first place. If ideologically misaligned individuals anticipate hostile institutional environments, selection into teaching or other public-facing professions subject to similar cultural policies may shift in ways that compound the sorting effects we show among existing workers.

References

Abrams, S. J. and Fiorina, M. P. (2012). "the big sort" that wasn't: A skeptical reexamination. *PS: Political Science & Politics*, 45(2):203–210.

Alexander, T. N., Clark, L. B., Flores-Ganley, I., Harris, C., Kohli, J., McLelland, L., Moody, P., Powell, N., Smith, M., and Zatz, N. (2024). CRT Forward Tracking Project. Technical report, UCLA School of Law Critical Race Studies Program.

Altman, S. (2025). Hiring and firing based on political views. *Southern California Law Review*, 99:1.

Anderson, L. (2010). Embedded, emboldened, and (net) working for change: Support-seeking and teacher agency in urban, high-needs schools. *Harvard Educational Review*, 80(4):541–573.

Ashforth, B. (2000). *Role transitions in organizational life: An identity-based perspective*. Routledge.

Ashforth, B. E. and Schinoff, B. S. (2016). Identity under construction: How individuals come to define themselves in organizations. *Annual Review of Organizational Psychology and Organizational Behavior*, 3:111–137.

Ashraf, N. and Bandiera, O. (2018). Social incentives in organizations. *Annual Review of Economics*, 10(1):439–463.

Bandiera, O., Best, M. C., Khan, A. Q., and Prat, A. (2021). The allocation of authority in organizations: A field experiment with bureaucrats. *The Quarterly Journal of Economics*, 136(4):2195–2242.

Barrett, N., Strunk, K. O., and Lincove, J. (2021). When tenure ends: The short-run effects of the elimination of Louisiana's teacher employment protections on teacher exit and retirement. *Education Economics*, 29(6):559–579.

Berkman, M. and Plutzer, E. (2010). *Evolution, creationism, and the battle to control America's* *classrooms*. Cambridge University Press.

Besharov, M. L. (2008). *Mission goes corporate: Employee behavior in a mission-driven business*. Harvard University.

Besley, T. and Ghatak, M. (2005). Competition and incentives with motivated agents. *American Economic Review*, 95(3):616–636.

Biasi, B. (2021). The labor market for teachers under different pay schemes. *American Economic Journal: Economic Policy*, 13(3):63–102.

Bishop, B. and Cushing, R. G. (2009). *The big sort: Why the clustering of like-minded America is tearing us apart*. Houghton Mifflin Harcourt.

Bissell, T. (2023). Teaching in the upside down: What anti-critical race theory laws tell us about the First Amendment. *Stanford Law Review*, 75:205.

Bolton, A., de Figueiredo, J. M., and Lewis, D. E. (2020). Elections, ideology, and turnover in the US federal government. *Journal of Public Administration Research and Theory*, 31(2):451–466.

Bonilla, S., Dee, T. S., and Penner, E. K. (2021). Ethnic studies increases longer-run academic engagement and attainment. *Proceedings of the National Academy of Sciences*, 118(37):e2026386118.

Broadwater, L. and Bender, M. C. (2026). Trump budget proposal targets "woke" programs. *The New York Times*. Accessed July 23, 2026.

Brown, J. R. and Enos, R. D. (2021). The measurement of partisan sorting for 180 million voters. *Nature Human Behaviour*, 5(8):998–1008.

Brown, J. R., Rendleman, H. E., and Silfa, M. (2026). How crisis reshapes government talent.

Brunner, E., Cowen, J. M., Strunk, K. O., and Drake, S. (2019). Teacher labor market responses to statewide reform: Evidence from Michigan. *Educational Evaluation and Policy Analysis*, 41(4):403–425.

Bunderson, J. S. and Thompson, J. A. (2009). The call of the wild: Zookeepers, callings, and the double-edged sword of deeply meaningful work. *Administrative Science Quarterly*, 54(1):32–57.

Carpenter, J. and Gong, E. (2016). Motivating agents: How much does the mission matter? *Journal of Labor Economics*, 34(1):211–236.

Čartolovni, A., Stolt, M., Scott, P. A., and Suhonen, R. (2021). Moral injury in healthcare professionals: A scoping review and discussion. *Nursing Ethics*, 28(5):590–602.

Chetty, R., Friedman, J. N., and Rockoff, J. E. (2014). Measuring the impacts of teachers II: Teacher value-added and student outcomes in adulthood. *American Economic Review*, 104(9):2633–2679.

Chinoy, S. and Koenen, M. (2024). Political sorting in the US labor market: Evidence and explanations.

Dewey, J. (1916). *Democracy and education: An introduction to the philosophy of education*. Macmillan, New York.

DiBenigno, J. (2022). How idealized professional identities can persist through client interactions. *Administrative Science Quarterly*, 67(3):865–912.

Dunn, A. H. (2018). Leaving a profession after it's left you: Teachers' public resignation letters as resistance amidst neoliberalism. *Teachers College Record*, 120(9):1–34.

Dunning, T. (2012). *Natural experiments in the social sciences: A design-based approach*. Cambridge University Press.

Duttmann, A. G. (2000). *Between cultures: Tensions in the struggle for recognition*. Verso.

Enamorado, T., Fifield, B., and Imai, K. (2019). Using a probabilistic model to assist merging of large-scale administrative records. *American Political Science Review*, 113(2):353–371.

Forand, J. G., Ujhelyi, G., and Ting, M. M. (2023). Bureaucrats and policies in equilibrium administrations. *Journal of the European Economic Association*, 21(3):815–863.

Frake, J., Hurst, R., and Kagan, M. (2026). Political segregation in the US workplace. *Nature*

Human Behaviour, pages 1–19.

Freidson, E. (2001). *Professionalism, the third logic: On the practice of knowledge*. University of Chicago Press.

Gailmard, L. and Gailmard, S. (2026). The persistence and fragility of bureaucratic capacity. Working Paper.

Gailmard, S. and Patty, J. W. (2007). Slackers and zealots: Civil service, policy discretion, and bureaucratic expertise. *American Journal of Political Science*, 51(4):873–889.

Geerlings, J., Thijs, J., and Verkuyten, M. (2019). Preaching and practicing multicultural education: Predicting students' outgroup attitudes from perceived teacher norms and perceived teacher–classmate relations. *Journal of School Psychology*, 75:89–103.

Gjefsen, H. M. and Gunnes, T. (2020). How school accountability affects teacher mobility and teacher sorting. *Education Economics*, 28(5):455–473.

Goldstein, D. (2024). Why it's hard to control what gets taught in public schools. *The New York Times*. Accessed: 2025-08-20.

Goncalves, M. S. O., Langrock, I., LaViolette, J., and Spoon, K. (2024). Book bans in political context: Evidence from US schools. *PNAS Nexus*, 3(6):pgae197.

Gusfield, J. R. (1986). *Symbolic crusade: Status politics and the American temperance movement*. University of Illinois Press.

Gutmann, A. (1999). *Democratic education: Revised edition*. Princeton University Press.

Haas, N. and Lindstam, E. (2024). My history or our history? Historical revisionism and entitlement to lead. *American Political Science Review*, 118(4):1778–1802.

Hackman, J. R. and Lawler, E. E. (1971). Employee reactions to job characteristics. *Journal of Applied Psychology*, 55(3):259.

Hall, M. and Hampden-Thompson, G. (2022). The teacher as street-level bureaucrat: Science teacher's discretionary decision-making in a time of reform. *International Journal of Science*

Education, 44(6):980–999.

Hayes, M. S. and Stazyk, E. C. (2019). Mission congruence: To agree or not to agree, and its implications for public employee turnover. *Public Personnel Management*, 48(4):513–534.

Hirschman, A. O. (1972). *Exit, voice, and loyalty: Responses to decline in firms, organizations, and states*. Harvard University Press.

Holman, M., Johnson, R., and Simko, T. (2025). Measuring conflict in local politics. *Urban Affairs Review*, page 10780874251355893.

Honig, D. (2024). *Mission driven bureaucrats: Empowering people to help government do better*. Oxford University Press.

Houston, D. M. (2024). Polarization, partisan sorting, and the politics of education. *American Educational Research Journal*, 61(3):508–540.

Hunter, J. D. (1992). *Culture wars: The struggle to control the family, art, education, law, and politics in America*. Avalon Publishing.

Ibarra, H. (1999). Provisional selves: Experimenting with image and identity in professional adaptation. *Administrative Science Quarterly*, 44(4):764–791.

Ingersoll, R. M. (2001). Teacher turnover and teacher shortages: An organizational analysis. *American Educational Research Journal*, 38(3):499–534.

Jackson, C. K. (2009). Student demographics, teacher sorting, and teacher quality: Evidence from the end of school desegregation. *Journal of Labor Economics*, 27(2):213–256.

Jacoby, W. G. (2014). Is there a culture war? Conflicting value structures in American public opinion. *American Political Science Review*, 108(4):754–771.

Jaffe, G., Schmitt, E., and Cooper, H. (2025). Hegseth is purging military leaders with little explanation. *The New York Times*. Accessed August 24, 2026.

Jeong, D. W. and Luschei, T. F. (2019). Teacher sorting within and across schools and nations: A comparative study of South Korea and the United States. *Teachers College Record*, 121(8):1–40.

Johnson, S. M. (2004). *Finders and keepers: Helping new teachers survive and thrive in our schools*. The Next Generation of Teachers. Jossey-Bass, San Francisco, CA.

Jones, M. and Hartney, M. T. (2017). Show who the money? Teacher sorting patterns and performance pay across US school districts. *Public Administration Review*, 77(6):919–931.

Jones, R. S. (1971). Teachers as agents of political socialization. *Education and Urban Society*, 4(1):99–114.

Khan, M. Y. (2025). Mission motivation and public sector performance: Experimental evidence from Pakistan. *American Economic Review*, 115(7):2343–2375.

Koedel, C., Podgursky, M., and Shi, S. (2013). Teacher pension systems, the composition of the teaching workforce, and teacher quality. *Journal of Policy Analysis and Management*, 32(3):574–596.

Kogan, V. (2022). Adult culture wars and student academic achievement. *Available at SSRN 4099637*.

Kogan, V. (2025). *No adult left behind: How politics hijacks education policy and hurts kids*. Cambridge Studies in the Comparative Politics of Education. Cambridge University Press.

Kraft, M. A. and Lyon, M. A. (2024). The rise and fall of the teaching profession: Prestige, interest, preparation, and satisfaction over the last half century. *American Educational Research Journal*, 61(6):1192–1236.

Kraft, M. A., Marinell, W. H., and Shen-Wei Yee, D. (2016). School organizational contexts, teacher turnover, and student achievement: Evidence from panel data. *American Educational Research Journal*, 53(5):1411–1449.

Larson, M. S. (1977). The rise of professionalism: A sociological analysis. *Class: The Anthology*, pages 263–286.

Leege, D. C., Wald, K. D., Krueger, B. S., and Mueller, P. D. (2009). *The politics of cultural* *differences: Social change and voter mobilization strategies in the post-New Deal period*. Princeton University Press.

Lipsky, M. (1980). *Street-level bureaucracy: Dilemmas of the individual in public service*. Russell Sage Foundation.

Lopes, J. and Oliveira, C. (2020). Teacher and school determinants of teacher job satisfaction: A multilevel analysis. *School Effectiveness and School Improvement*, 31(4):641–659.

Lyon, M., Snaider, C., Bertrand, A., Bak, J., and Jacobsen, R. (2026). Racial status threat and anti-DEI efforts in local school districts. Working Paper.

Maitlis, S. and Christianson, M. (2014). Sensemaking in organizations: Taking stock and moving forward. *Academy of Management Annals*, 8(1):57–125.

Martin, G. J. and Webster, S. W. (2020). Does residential sorting explain geographic polarization? *Political Science Research and Methods*, 8(2):215–231.

Maynard-Moody, S. W. and Musheno, M. C. (2003). *Cops, teachers, counselors: Stories from the front lines of public service*. University of Michigan Press.

McNeil, E. (2015). Survey explores why people go into teaching in the first place. *Education Week*. Accessed: 2026-06-27.

Mehta, J. (2013). *The allure of order: High hopes, dashed expectations, and the troubled quest to remake American schooling*. Oxford University Press.

Moe, T. M. (2011). *Special interest: Teachers unions and America's public schools*. Bloomsbury Publishing USA.

Moe, T. M. (2019). *The politics of institutional reform: Katrina, education, and the second face of power*. Cambridge University Press.

Mummolo, J. and Nall, C. (2017). Why partisans do not sort: The constraints on political segregation. *The Journal of Politics*, 79(1):45–59.

Nalewajko, K. (2025). Allies of the weak: La Résistance and Jews in the Holocaust. *American*

Political Science Review, 119(4):1942–1962.

National Center for Education Statistics (2024). Teacher turnover: Stayers, movers, and leavers. Technical report, U.S. Department of Education, Institute of Education Sciences. Accessed July 23, 2026.

National Education Association (2023). Teacher tenure & due process protections for educators. National Education Association. Accessed July 23, 2026.

Nelsen, M. D. (2023). *The color of civics: Civic education for a multiracial democracy*. Oxford University Press.

Nguyen, T. D., Pham, L. D., Crouch, M., and Springer, M. G. (2020). The correlates of teacher turnover: An updated and expanded meta-analysis of the literature. *Educational Research Review*, 31:100355.

Oldmixon, E. A. (2002). Culture wars in the congressional theater: How the US House of Representatives legislates morality, 1993–1998. *Social Science Quarterly*, 83(3):775–788.

Paglayan, A. (2026a). When does education promote political participation? The role of curriculum policies. *Working Paper*.

Paglayan, A., Neundorf, A., and Kim, W. (2026). Democracy and the politicization of education. Available at SSRN 4395806.

Paglayan, A. S. (2026b). Education as a political institution. *Annual Review of Political Science*, 29(1):525–546.

Patel, V. (2026). Professors are being watched: 'We've never seen this much surveillance'. *The New York Times*.

Pawlewicz, D. D. A. (2020). *Blaming teachers: Professionalization policies and the failure of reform in American history*. Rutgers University Press.

PEN America (2024). PEN America index of school book bans.

Perry, J. L. and Wise, L. R. (1990). The motivational bases of public service. *Public Adminis-*

tration Review, pages 367–373.

Peterson, E., Simko, T., Joyce, M., McCrain, J., and White, R. (2025). The personnel costs of nationalization: Culture war conflicts lead to higher superintendent turnover. Working Paper.

Pierskalla, J. H. and Sacks, A. (2020). Personnel politics: Elections, clientelistic competition and teacher hiring in Indonesia. *British Journal of Political Science*, 50(4):1283–1305.

Pignatelli, F. (1993). What can I do? Foucault on freedom and the question of teacher agency. *Educational Theory*, 43(4):411–432.

Pollock, M., Rogers, J., Kwako, A., Matschiner, A., Kendall, R., Bingener, C., Reece, E., Kennedy, B. C., and Howard, J. R. (2022). The conflict campaign: Exploring local experiences of the campaign to ban "critical race theory" in public K–12 education in the US, 2020–2021. Technical report.

Pratt, M. G. (2000). The good, the bad, and the ambivalent: Managing identification among Amway distributors. *Administrative Science Quarterly*, 45(3):456–493.

Pratt, M. G., Rockmann, K. W., and Kaufmann, J. B. (2006). Constructing professional identity: The role of work and identity learning cycles in the customization of identity among medical residents. *Academy of Management Journal*, 49(2):235–262.

Rahnama, R. (2025). Ideology as compensation: When and how dominant groups deploy ideology to sustain hierarchy. *Available at SSRN 5297720*.

Ravitch, D. (1990). Multiculturalism: E pluribus plures. *The American Scholar*, 59(3):337–354.

Renkl, M. (2022). A thank-you note to teachers after a year of attacks. *The New York Times*. Accessed July 23, 2026.

Richardson, M. D. (2019). Politicization and expertise: Exit, effort, and investment. *The Journal of Politics*, 81(3):878–891.

Rofes, E. (2005). *A radical rethinking of sexuality and schooling: Status quo or status queer?* Bloomsbury Publishing PLC.

Rosso, B. D., Dekas, K. H., and Wrzesniewski, A. (2010). On the meaning of work: A theoretical integration and review. *Research in Organizational Behavior*, 30:91–127.

Rovner, J. and Pradhan, R. (2024). More medical residents are avoiding states with abortion restrictions, analysis finds. *PBS NewsHour*.

Sapiro, V. (2004). Not your parents' political socialization: Introduction for a new generation. *Annual Review of Political Science*, 7(1):1–23.

Schabram, K. and Maitlis, S. (2017). Negotiating the challenges of a calling: Emotion and enacted sensemaking in animal shelter work. *Academy of Management Journal*, 60(2):584–609.

Schneider, B. (1987). The people make the place. *Personnel Psychology*, 40(3):437–453.

Seawright, J. (2016). *Multi-method social science: Combining qualitative and quantitative tools*. Cambridge University Press.

Shao, E. and Wu, A. (2025). The federal work force cuts so far, agency by agency. *The New York Times*. Accessed August 24, 2026.

Sinha, S., Zerbino, N., Valant, J., and Perera, R. M. (2023). Moms for Liberty: Where are they, and are they winning? Technical report, Brookings Institution.

Sorensen, L. C. and Ladd, H. F. (2020). The hidden costs of teacher turnover. *AERA Open*, 6(1):2332858420905812.

Spenkuch, J. L., Teso, E., and Xu, G. (2023). Ideology and performance in public organizations. *Econometrica*, 91(4):1171–1203.

Stryker, S. and Burke, P. J. (2000). The past, present, and future of an identity theory. *Social Psychology Quarterly*, pages 284–297.

Sutcher, L., Darling-Hammond, L., and Carver-Thomas, D. (2019). Understanding teacher shortages: An analysis of teacher supply and demand in the United States. *Education Policy Analysis Archives*, 27(35).

Tiebout, C. M. (1956). A pure theory of local expenditures. *Journal of Political Economy*, 64(5):416–424.

Torres, E. (2026). Americans' attitudes about political neutrality in public schools. EdWorkingPaper 26-1431, Annenberg Institute at Brown University.

Tummers, L. (2012). Policy alienation of public professionals: The construct and its measurement. *Public Administration Review*, 72(4):516–525.

Van Maanen, J. and Barley, S. R. (1982). Occupational communities: Culture and control in organizations. Technical report.

Vandenabeele, W. (2008). Government calling: Public service motivation as an element in selecting government as an employer of choice. *Public Administration*, 86(4):1089–1105.

Weatherley, R. and Lipsky, M. (1977). Street-level bureaucrats and institutional innovation: Implementing special-education reform. *Harvard Educational Review*, 47(2):171–197.

Weiss, D. C. (2025). Resignations multiply at BigLaw firms that made deals with Trump. *ABA Journal*.

Wildavsky, A. (1987). Choosing preferences by constructing institutions: A cultural theory of preference formation. *American Political Science Review*, 81(1):3–21.

Wilson, J. Q. (1989). *Bureaucracy: What government agencies do and why they do it*. Basic Books.

Wirsching, E. (2026a). Political power of bureaucratic agents: Evidence from policing in New York City. *The Journal of Politics*.

Wirsching, E. M. (2026b). The limits of merit: Career pipelines and representation in US law enforcement. Working Paper.

Woo, A., Diliberti, M. K., and Steiner, E. D. (2024). Policies restricting teaching about race and gender spill over into other states and localities. RAND Policy Report. https://www.rand.org/pubs/research_reports/RRA1108-10.html.

Worth, J. and Van den Brande, J. (2020). Teacher autonomy: How does it relate to job satisfaction and retention? *National Foundation for Educational Research*.

Wrzesniewski, A. and Dutton, J. E. (2001). Crafting a job: Revisioning employees as active crafters of their work. *Academy of Management Review*, 26(2):179–201.

Xu, Y., Zhao, A., and Ding, P. (2026). Factorial difference-in-differences. *Journal of the American Statistical Association*, pages 1–23. just-accepted.

Yair, O. and Sulitzeanu-Kenan, R. (2018). When do we care about political neutrality? The hypocritical nature of reaction to political bias. *PLoS ONE*, 13(5):e0196674.

Zacka, B. (2017). *When the state meets the street: Public service and moral agency*. Harvard University Press, Cambridge, MA.

Zimmerman, J. (2024). *Whose America?: Culture wars in the public schools*. University of Chicago Press.

Appendix

Contents

A Conflict Categorization Scheme

As described in the main text, we create several variables about each conflict in our dataset: (1) the specific employee roles that are most likely to be implicated, (2) the ideological orientation of the final decision, and (3) whether or not a policy was ultimately enacted or not. This section describes how these variables are coded for each dataset.

PEN Book Bans

- **Exposure:** We treat English/LA and librarians as most directly impacted.
- **Ideology:** We treat all enacted book bans as pro-Conservative.
- **Enacted:** PEN tracks whether each ban was ultimately enacted or not, which we adopt.

CRT Forward

- **Exposure:** We treat English/LA, librarians, and History/Social Studies teachers as most directly impacted.
- **Ideology:** We treat all enacted CRT policies and book bans as pro-Conservative, because manually coding the ideological direction of each proposal did not reveal any instances of pro-liberal policies. For example, a pro-liberal policy could require the instruction of CRT in classrooms. We did not identify any such policies in any of our datasets.
- **Enacted:** CRT Forward collects information on whether each policy was enacted or failed, which we adopt.

Cato Battle Map

- **Exposure:** Cato contains thousands of distinct conflicts that expose different roles. We pass the conflict "description," typically a paragraph of text describing the nature of the conflict and final outcome, to Claude Sonnet 4.6 to predict which roles from our teacher categorization scheme are most impacted by each conflict. For categorizing CATO observations into exposure, ideology, and enacted values, we use the prompt shown in Appendix A.1.
- **Ideology:** Similarly, we also use Sonnet 4.6 to predict whether the ideological lean of the final outcome was pro-Conservative, Liberal, or neither.
- **Enacted:** Finally, we also classify whether each conflict led to a policy change/adoption or not.

A.1. CATO Classification Prompt

You are a research assistant classifying K-12 culture-war conflicts for an academic study on teachers as street-level bureaucrats. You assess how a conflict bears on the professional discretion of classroom-level staff - teachers, counselors, librarians (NOT principals, district administrators, or other non-instructional staff) - in delivering their public work: what they may, must, or must not do in instruction, material selection, or student interaction. You will also assess the perceived ideology of the result.

NATURE {none, indirect, direct}. Resolve by working the two gates in order.

GATE 1 - Does the conflict touch classroom-staff work in any way - through their instruction, materials, student interaction, OR an enforcement duty placed on them?

- No -> none. The conflict places no obligation of any kind on classroom staff: it targets only students with no staff role (e.g., enrollment, athletic eligibility, busing, facilities, mascots); targets principals/district administrators; or concerns the private, off-duty life of a staff member. Other examples: a dispute over the appointment or conduct of a principal; a teacher reprimanded for personal social-media posts unrelated to teaching; a purely symbolic board resolution imposing no requirement on classroom staff.
- Yes -> go to Gate 2. Note: a student-targeted rule that staff must ENFORCE (e.g., a student dress code) or that affects student-teacher interactions passes this gate on the enforcement duty alone, and will resolve to indirect at Gate 2 - it is not none.

GATE 2 - Does it change the core exercise of professional judgment, or bear on it only at one remove?
Ask: does the conflict (or a measure resulting from it) change what or how staff teach, what materials they may use, or how they address/interact with students - applied as a general rule?

- Yes, general rule changing core discretion -> direct. Changes includes EXPANSION, not only restriction. Examples: a curriculum requirement or instructional ban (e.g., staff may not teach material related to race/racism); a teacher dress or conduct code; a policy governing how staff must address students (chosen pronouns, legal names); an AI-use policy governing teachers; a policy affirming staff latitude to discuss contested topics (expansion).
- No, only at one remove -> indirect. This holds when EITHER (a) the conflict primarily targets students but creates an enforcement role for staff without changing how they instruct or interact pedagogically - e.g., a student dress code; OR (b) it concerns the individual conduct of a single staff member rather than a general rule. Examples: an individual

teacher placed on leave for off-hand offensive language; a single teacher disciplined for teaching a challenged book or for refusing to use student pronoun requests.

STATUS {enacted, not_enacted}. Did the conflict result in a binding measure reaching classroom staff being adopted or revised? enacted = a districtwide (or higher) written rule was adopted or revised; not_enacted = proposed but failed, OR contestation that never produced a binding rule. ORTHOGONAL to NATURE - do not let enactment influence the nature label, or vice versa. Directness is determined by what a measure targets, NOT by whether it passed: a proposed-but-failed pronoun mandate is still direct in NATURE and not_enacted in STATUS.

PRIORITY RULES

1. Resolve NATURE to the most direct applicable label (direct > indirect > none). If a conflict bundles provisions (e.g., a curriculum mandate AND a student dress code), code the most direct provision that reaches classroom staff -> direct.

2. Enforcement spillover is NOT directness. Teachers having to enforce a student-targeted rule -> at most indirect. Directness (Gate 2 = yes) requires the rule to change instruction, materials, or pedagogical interaction itself.

3. The case of a single staff member is indirect, regardless of what the underlying conduct was, because it is not a general rule.

DECISION RULES FOR HARD CASES

- Student dress code -> indirect (enforcement spillover), never direct, never none. Teacher dress code -> direct.
- Pronoun/name-address rule for staff, as a general policy -> direct, whether or not enacted. A single teacher disciplined over pronouns with no general rule -> indirect.
- AI policy naming both students and teachers -> direct.
- Off-duty/private-life behavior with no instructional nexus -> none, even if the staff member is disciplined.
- Principal/administrator conduct -> none, even if teachers are affected downstream.
- A proposed measure that failed -> code NATURE on what it would have targeted; STATUS = not_enacted.

IDEOLOGY {pro-conservative, pro-liberal, neither}. The ideological direction of the final outcome:

- “pro-conservative” — restricts DEI discussions, introduces religious content, or otherwise moves in a conservative direction.
- “pro-liberal” — expands DEI discussions, removes religious content, or otherwise moves in a liberal direction.
- “neither” — unclear or not applicable.

For ideology, code the direction of the most recent concluded action. If a challenge or proposed policy was rejected or failed, code the direction of the outcome (e.g., a failed attempt to remove a book means books stay -> “pro-liberal”). If no action has concluded, code the direction of the original proposal.

EXPOSED_STAFF — Which classroom-level staff roles would be most directly and concretely affected by this conflict or its outcome? Return a JSON array. If nature is “none”, return []. Multiple values are encouraged when a conflict spans roles.

Role definitions:

- “all” — use only when the policy applies uniformly across all subject areas and classroom types by its own terms. Use for: general staff conduct or dress codes; pronoun or preferred-name mandates directed at all staff; classroom display or decor mandates (flags, posters, signs); Pledge of Allegiance policies; mandatory professional development or training requirements. Do NOT use “all” simply because the subject matter is hard to classify — see the subject-unclear rule below.
- “librarian” — book challenges, library collection policies, access restrictions, or removal of materials from library shelves. Use even when the challenged book also appears on a course reading list; in that case, also add “english_la”.
- “english_la” — conflicts over novels, literary texts, poems, or reading lists assigned in English or language arts courses. Add “librarian” if library access is also at issue.
- “history_social_studies” — conflicts over history, civics, patriotism, ethnic studies, social studies curricula, or race/racism in a non-literary instructional context (e.g., CRT bans, anti-racist curriculum, Pledge of Allegiance content disputes). Do not use for book challenges where the primary venue is the library or English class.
- “health_pe” — sex education, reproductive health, contraception, gender in athletics, bathroom or locker room access policies that create an enforcement or instructional role for staff, or physical education policies.

- “science” — evolution, creationism, intelligent design, climate change, fossil fuels, biology textbooks, or other science-content disputes.
- “arts” — theater productions, play selection or cancellation, music, choir, band, visual art, or film/media courses. Use when the conflict centers on content delivered by arts or performing-arts staff.
- “counselor” — school counselor-specific curricula (e.g., social-emotional learning programs assigned to counselors), mental health referral protocols, or guidance policies that govern how counselors interact with students.
- “coach_extracurricular” — conflicts primarily affecting coaches or extracurricular staff: prayer or religious expression at sporting events, student political expression at games or in uniform, eligibility disputes that create a role for coaches, or after-school club policies.
- “math” — math textbook or math curriculum content disputes (rare; use only when the conflict explicitly concerns mathematics instruction).
- “esl_bilingual” — language-of-instruction mandates, bilingual education policies, or English-learner program disputes.
- “special_ed” — conflicts that directly implicate special education teachers or IEP (individualized education plan)-related staff: parental rights disputes over IEP content or process; inclusion vs. self-contained placement policies that create instructional obligations; disability accommodations that interact with culture-war topics (e.g., gender identity written into an IEP, parental-notification rules that conflict with IEP-mandated supports); or behavioral/disciplinary policies that fall disproportionately on SPED staff.
- “other” — classroom-level staff are concretely affected but no specific role above fits.

Pairing guidance:

- Humanities, DEI, anti-racist, or CRT curricula with no single subject specified -> [“english_la”, “history_social_studies”].
- A challenged book that appears in both the library and on a course reading list -> [“librarian”, “english_la”].
- Prayer or religious expression at school events -> [“coach_extracurricular”]; if it extends into classroom instruction or all-staff conduct, also add “all”.

- A pronoun or preferred-name policy directed at all staff -> ["all"]; a single teacher disciplined over pronouns with no general policy -> ["all"] still applies for exposure (the teacher is in a general classroom role), but nature will be indirect.
- Gender identity or parental-notification conflicts where the accommodation is written into an IEP or specifically involves a student with a disability -> add "special_ed" alongside any other applicable role (e.g., ["special_ed", "all"]).

Subject-unclear rule: If staff are concretely affected but the description does not state or clearly imply which subject area or staff role is involved, return ["all"] rather than inferring a subject. Do not infer a role from topic alone — a district-wide assignment on systemic racism with no named class -> ["all"], not ["history_social_studies"].

Output format:
Respond ONLY with a JSON object with exactly four fields: nature, status, ideology, exposed_staff. No explanation, no markdown, no backticks.

A.2. Manual Validation of a Subset of CATO Conflicts

We validate Claude's classification of CATO incidents against independent human coding. From the 4,345 eligible coded incidents, we draw a reproducible random sample of 400 (roughly 9% of the corpus). A research assistant blindly coded each sampled incident's title and description text using the same codebook given to Claude. We then compare the human codes to Claude's original classification of the same 400 incidents on all four coding dimensions: *nature* (whether the conflict bears directly, indirectly, or not at all on classroom-staff discretion), *status* (enacted vs. not enacted), *ideology* (pro-conservative, pro-liberal, or neither), and *exposed staff* (which classroom roles the conflict most directly implicates; incidents may carry multiple role tags).

For a given label (e.g., ideology = pro-conservative), precision is the share of incidents Claude assigned that label where the RA agreed, and recall is the share of incidents the RA assigned that label where Claude agreed. Table A1 reports exact agreement for each dimension, together with precision and recall averaged across that dimension's labels (macro-averaged for nature/status/ideology; micro-averaged across all role tags for exposed staff, since an incident can carry more than one).

Classifier reliability varies by dimension. Agreement is strongest for status and nature (0.68 and 0.64 exact agreement, respectively), where precision and recall are also reasonably balanced. Ideology is the weakest dimension: exact agreement is 53%, and precision on the two partisan

Table A1. CATO Classification Validation: Claude vs. Blind Human Coding (N = 400)

Dimension	Exact agreement	Precision	Recall
Nature (none / indirect / direct)	0.64	0.62	0.64
Status (enacted / not enacted)	0.68	0.67	0.75
Ideology (pro-con. / pro-lib. / neither)	0.53	0.56	0.58
Exposed staff (role tags, multi-label)	0.53	0.63	0.60

Notes: Precision and recall are macro-averaged across a dimension's labels, except exposed staff, which is micro-averaged across role tags because incidents may carry more than one.

labels is low in the full breakdown (0.41 for both pro-conservative and pro-liberal), meaning that when Claude assigns either partisan label, the blind human coder agrees less than half the time. Status precision is similarly uneven across its two values — recall on "enacted" is high (0.88) but precision is low (0.39), indicating Claude over-assigns that label relative to human coding. Because CATO incidents are one of three sources feeding our treatment classification, and ideology in particular is the dimension most analogous to the alignment construct central to our argument, we treat CATO-based ideology and enactment-status codings as noisier than the CRT and book-ban treatment sources, and flag this as a limitation of the CATO measure.

Table A2. Illustrative Conflict Examples from our Dataset. Each row represents one documented conflict from the CATO dataset. *Exposed Staff* identifies the public school employee roles most directly implicated according to our identification process described in the main text. *Ideology* indicates the predicted ideological lean of the final result. *Outcome* indicates whether the initiating side's position was ultimately adopted (**Enacted**) or rejected / stalled (**Not Enacted**).

Conflict Description	**District, State (Year)**	**Exposed Staff**	**Ideology**	**Outcome**
Parents raised concerns about sexually explicit content in the district's sex ed curriculum. The curriculum provider, Planned Parenthood, restricted access to enrolled students' parents only. Citing a rule requiring materials be available for public inspection, the superintendent dropped the course. A teacher supporter said, "I feel it's been very valuable, not only for my children, but for the children of the entire community."	Pacific Grove USD, California (2024)	Health/PE	Con.	**Enacted**
Marian: The True Story of Robin Hood, a play featuring LGBTQ+ characters, was canceled after parents complained the subject matter was inappropriate. Supporters petitioned to reinstate it, writing, "We need to show students that there is nothing wrong with being who they are unapologetically." Officials cited "safety concerns."	Northwest Allen County SD, Indiana (2023)	Arts	Con.	**Enacted**
[The] district removed 150 books for review following parent complaints. The district's curriculum coordinator said, "Some people feel we're not doing enough, some feel we're going too far… We all want to ensure that students have access to worthy reading materials."	Coeur d'Alene SD, Idaho (2025)	Librarian; English/LA	Con.	**Enacted**
Ready Player One, used in an English class, was challenged at a board meeting. The challenger cited "graphic sexual content, foul language, and statements that compare faith in God to belief in the Easter Bunny." A parent supporter said the book made her son want to read for the first time in years. The board voted to keep the book.	North Gibson SD, Indiana (2025)	English/LA	Liberal	**Not Enacted**
A board member alleged that a middle school civics lesson contained a biased portrayal of the Republican National Committee's energy platform. An opposing board member said any curriculum concern "would be something for the principal to address with the staff member." The board voted against changing the curriculum.	Warren County SD, Virginia (2025)	History / Social Studies	Liberal	**Not Enacted**

B Classifying Staff Roles and Subject Exposure

B.1. Motivation

The administrative rosters obtained from our sample states record staff roles and teaching subjects in distinct ways. Some states pack both pieces of information into a single free-text job title (Florida, Illinois), some split them across a role field and a subject field (Kansas, Iowa), and some entangle several roles and subjects within one record (Kentucky, Idaho). To analyze turnover by position and by exposure to the subjects targeted by a culturally conflictual policy, we map this heterogeneous text onto a common set of binary indicators that is identical across states.

We classify staff with a large language model. Each unique role/subject descriptor is assigned a set of binary flags, which are then merged back onto every roster row that shares the descriptor. The classification was produced with Gemini 2.5 Flash through its batch interface; the pipeline is model-agnostic and was validated against an alternative run with Claude. Because classification operates on the set of unique descriptors rather than on individual rows, the cost is governed by the number of distinct titles (a few hundred to a few thousand per state) rather than by the millions of staff-years.

B.2. Procedure

For each roster row, we concatenate every column that carries role or subject information into a single descriptor string, with segments separated by a delimiter. This sidesteps the cross-state column problem: a counselor identified through a subject field in one state and through a position field in another is captured either way, because the model sees all of the available text. Where a state distinguishes certified from classified staff (Georgia, Kentucky, Idaho, Iowa), a readable prefix recording that status is prepended to the descriptor and added to the classification key, so that the same title held by certified and by classified staff is classified separately. Teaching certifications and endorsements are deliberately excluded from the input, because they record what a staff member is licensed to teach rather than what they actually teach, and would bias the subject indicators toward licensed fields.

The model returns, for each descriptor, a multi-label assignment: any number of the categories may apply, since a single record can name several roles and several subjects. Each category is an independent indicator taking the value 1 (applies), 0 (assessed and does not apply), or missing.

B.3. Categories

We distinguish ten *position* indicators and eight *subject* indicators (Tables A3 and A4). The four role-based categories (librarian, counselor, coach, special education) are treated as position indicators read from the role text, and are thus observed (never missing) in every state. The specific role flags and the residual support-staff category are mutually carved so they do not double-count, while distinct roles can co-occur (e.g., a special-education teacher is flagged both as a teacher and as special education).

Table A3. Position indicators (ten, always observed)

Indicator	Covers
`teacher`	classroom teacher or instructor
`superintendent`	district superintendent
`admin`	principal, assistant superintendent, administrative/clerical leadership
`librarian`	librarians and library media specialists
`counselor`	guidance counselors, school psychologists, social workers, career advisors
`coach_extracurricular`	athletic coaches and extracurricular advisors
`special_ed`	special-education teachers, exceptional-child educators, aides serving students with disabilities
`support_staff`	other support roles (paraprofessionals, instructional coaches, clerical, custodial, food service, transport, nurse)
`substitute`	substitute or temporary teaching positions
`other_position`	none of the above

Table A4. Subject-exposure indicators (eight)

Indicator	Covers
`english_la`	English, reading, writing, literacy, composition
`history_social_studies`	history, social studies, civics, government, ethnic studies
`health_pe`	health, physical education, sex education, health sciences
`science`	biology, chemistry, physics, earth science, engineering, computer science
`arts`	fine and performing arts, music, drama, dance
`math`	mathematics, statistics, data science
`esl_bilingual`	world languages other than English, English as a second language
`other_subject`	instructional staff teaching some other identifiable field

B.4. Treatment of missing subject information

The subject indicators distinguish a true zero from missing information, because "we have no subject data for this row" must not be read as "assessed and not this subject." When at least one teaching subject is identifiable on a record, the corresponding indicators are set to 1 and the remaining subject indicators are set to 0. When no teaching subject is identifiable, which is the case for non-instructional roles and for generic teaching titles that name a grade level but no subject, every subject indicator is set to missing. As a result, a record contributes a 0 to a subject only when it has been positively assessed as teaching some other subject.

Two states, Georgia and Kentucky, record no academic-subject field at all. For these states the role-based indicators remain fully identified from the role text, and a small number of specialty subjects appear in specific titles, but the core academic subjects (English/language arts, history/social studies, science, mathematics) are not identifiable for ordinary classroom teachers and are returned as missing.

B.5. Per-state inputs

Table A5 lists the columns concatenated into the descriptor for each state. Kentucky and Idaho records concatenate several roles, so multi-role assignments are common; Kansas, Iowa, Texas, Connecticut, Wisconsin, and Missouri records can name several subjects; Florida and Illinois embed the subject within the job title; Pennsylvania's assignment text carries role, subject, and grade range together. The number of distinct descriptors, which governs the cost of classification, ranges from a few dozen (Illinois, Florida) to tens of thousands (Idaho, Iowa), where several assignments are concatenated within a single record.

B.6. Validation and limitations

Each state was first run on a pilot sample and the assigned flags inspected before the full classification. For states carrying a certified/classified distinction, a cross-tabulation of the teacher indicator against that status confirms that classified staff are almost never labeled teachers. The resulting per-state lookups share a constant eighteen-indicator schema, so they stack for our pooled analyses. The primary limitation is that subject exposure is only as rich as each state's text: in the two states without an academic-subject field, exposure to the core academic subjects cannot be measured for ordinary classroom teachers and is treated as missing rather than as absence of exposure.

Table A5. Role and subject text used for classification, by state

State	Columns concatenated	Descriptors	Academic subject present
FL	job title	125	yes (in title)
TX	job title + subject-area list	3,366	yes
GA	job title	577	no
CT	job-classification type	146	yes
IL	position description	44	partial
KS	educator type + subject-area list	2,917	yes
IA	position name + assignment names	23,145	yes
KY	job-class description	6,723	no
ID	assignment codes	17,501	yes (role and subject entangled)
PA	position + assignment description	385	yes
WI	assignment position + area + staff category	308	yes
MO	position title + course name	2,089	yes

C Linking Employees to the Voter Files

C.1. Objective and data

To observe employees' party of registration, we link the administrative rosters to the L2 national voter file. The voter side is a deduplicated unique-voter file carrying, for each registered voter, first/middle/last name, date of birth, sex, county, and party information. Both sides are subset to the relevant state before linkage. The match is probabilistic, because no common identifier exists between the rosters and the voter file.

C.2. Method

We use the Fellegi–Sunter probabilistic record-linkage model as implemented in `fastLink`, with Jaro–Winkler string comparison for names. Because enumerating all employee-by-voter pairs is infeasible at our large scale, we generate candidate match pairs by multi-pass blocking and then score the candidate set with a single model.

Blocking. Candidate pairs must share a first-name blocking key and, where the data allow, agree on sex and fall within one year on date of birth. The first-name key is the union of two encodings, an alphabetic k-means cluster and a Double Metaphone phonetic code, so that spelling and phonetic variants of a given name (for example Catherine and Katherine, Sean and Shawn) are both retained. We do not block on surname to avoid under-matching working-age women who change their surnames over time. Observations missing sex or date of birth are not immediately discarded, so that they could possibly be matched.

The fields available for blocking differ by state (Table A6). Five states have usable date of birth and sex; five carry sex only; two carry neither, so for those the candidate set rests on the name key alone.

Table A6. Identity fields available for linkage, by state

Regime	States	Blocking keys
date of birth + sex	FL, TX, GA, CT, IA	name $\times$ sex $\times$ DOB($\pm$1)
sex only	KY, KS, PA, WI, MO	name $\times$ sex
neither	IL, ID	name only

Florida and Connecticut have an exact date of birth; Texas, Georgia, and Iowa derive it from reported age and are therefore accurate only to within about a year, which the $\pm$1-year blocking window accommodates. The middle initial, where present, is scored rather than blocked on (it is too sparse in some states to serve as a blocking key), and is reported alongside the match rates in Table A7.

Scoring with one global model. Rather than estimating a separate model within each block, we fit a single Fellegi–Sunter model by the EM algorithm to the union of all candidate pairs. This choice is dictated by the structure of the candidate set: name blocking produces tens of thousands of small blocks (a median on the order of ten candidate pairs each). This is far too small to estimate a stable model block by block, whereas a single model fit on the pooled candidate set is fully powered and places every pair on one comparable probability scale, which makes a single state-wide acceptance threshold coherent. Fitting on the candidate union preserves the concentration of true matches that makes the parameters identifiable in a setting where the population match rate is on the order of 10^{-7}. The implementation reproduces `fastLink`'s agreement classification and uses its EM estimator; we departs from the default package implementation only in fitting a single pooled model, in using a vectorized blocker and agreement counter (for tractability), and in a custom date-of-birth comparator.

Scored fields and missingness. We score first name, last name, and, when possible, middle initial and date of birth. Sex is used for blocking only, since exact agreement within a block leaves it no discriminating variation. Date of birth is scored in addition to its blocking role, because the $\pm$1 window leaves residual exact versus off-by-one variation that breaks ties among voters who share a name, sex, and birth year. A field missing on either side is treated as missing-at-random and marginalized out of that pair's likelihood, following Enamorado et al.

(2019), so it is never read as a disagreement. A field is dropped from the model only when it is missing for every record in a state.

Acceptance, geography, and resolution. Pairs with posterior match probability above a threshold of 0.85 are retained. A soft geographic filter is applied after thresholding rather than as a block: a matched pair is kept only if the voter's county lies within 150 miles of a district in which the employee worked. Each employee is matched as a person rather than as a employee-by-year record, and is resolved to at most one voter (the highest-probability match), while the voter side retains all of its name, county, and birth-date variants so that a employee recorded under a maiden or married name can still match. Employees tied at the maximum posterior are held out as unresolved.

C.3. Match rates

Table A7 reports match rates for all twelve linked states. The rate is the share of distinct employees successfully matched to a voter. Only states carrying a date of birth matches between 77 and 86 percent: the birth date is what separates voters who share a name and sex, and without it a large share of employees tie at the maximum posterior and are held out rather than guessed.

Among the states without a date of birth, middle initials help to improve match rates. Kansas and Idaho, which carry an initial for roughly four employees in five, reach 75 and 74 percent; Idaho achieves this without even a sex field. At the other extreme, Wisconsin and Missouri have effectively no middle initial and sit at 36 and 44 percent, the lowest of the twelve. Between them, Pennsylvania and Illinois carry an initial for roughly two employees in three and land near 59 percent.

We also find a relatively low match in rates in Kentucky and Wisconsin, which appear to have many out-dated registrations on their voter file. Kentucky has 5.4 million distinct registered voters for their population of roughly 4.5 million, and Wisconsin similarly has 7.1 million entries on the voter file for a population of 5.9 million. While these records likely consist of inactive, moved, and deceased registrations, they likely lower our match rate by creating ambiguity. Without a birth date, two voters who share a name and sex produce the same agreement pattern and are assigned the same posterior, so a employee whose name is carried by several registrations ties at the maximum and is held out.

Table A7. Employee-to-voter match rates

State	Employees	Match rate	Identity fields (middle initial coverage)
GA	412,089	85.7%	age-derived DOB + sex + middle (80%)
CT	73,934	85.5%	exact DOB + sex + middle (85%)
TX	847,542	83.9%	age-derived DOB + sex
FL	294,189	77.2%	exact DOB + sex
IA	139,591	77.2%	age-derived DOB + sex + middle (7%)
KS	60,321	75.3%	sex + middle (81%)
ID	67,625	73.7%	middle (83%)
PA	195,418	59.6%	sex + middle (63%)
IL	246,298	58.9%	middle (71%)
KY	183,339	50.6%	sex + middle (82%)
MO	162,060	44.0%	sex
WI	224,160	36.1%	sex

C.4. Party of registration

Having linked employees to voters, we attach party from the voter file. The meaning of the party variable differs by state according to whether the state registers voters by party, and the two kinds of measure are not directly comparable. In party-registration states (Florida, Pennsylvania, Connecticut, Kentucky, Kansas, Iowa, Idaho) the full registration label is observed, including minor parties, and a change over time reflects an actual re-registration. In states without party registration (Texas, Georgia, Illinois, Wisconsin, Missouri) no registration exists to observe, and party is instead inferred from which primary ballot a voter took. A change then reflects a switch in the primary participated in rather than any change in affiliation. For the analysis, we collapse party into Democratic, Republican, and other, and use the party recorded at the time a employee is first treated, held fixed thereafter.

C.5. Limitations

We note three limitations of our merging approach. First, coverage rates vary across states as a function of their data availability. If match failures are correlated with features like ethnicity and sex, our downstream estimates conditioning on party would be compromised (see below for how we account for this limitation). Second, the procedure is deliberately conservative at ties: a employee who cannot be separated from several same-name voters is dropped rather than assigned, which protects precision at the cost of recall and drives the low rates in Wisconsin and Missouri. Third, in states whose employee identifier is constructed from name and birth date, distinct individuals who share both are merged, and in states with no birth date at all the identifier rests on the name alone.

C.6. Addressing representativeness of merged sample

The first limitation above can be measured directly. For each state we compare the demographics recorded in the administrative roster between public school employees who were linked to a voter and those who were not, using whatever fields the state collects (race, sex, salary, experience, degree, certification, and grade level; the set differs by state, and Illinois and Idaho record no race or sex at all). We summarise each comparison with a standardized mean difference: the gap in group means in pooled standard-deviation units, which is comparable across variables measured on different scales and, unlike a significance test, does not simply reflect the very large sample sizes.

We find some demographic patterns in match rates. Across all 205 field-by-state comparisons, 19 percent exceed the conventional $|\text{SMD}| = 0.10$ threshold for meaningful imbalance, though none are large in Cohen's sense (all are below 0.5). Merged employees are more likely to be White and less likely to be Hispanic or Asian, more likely to be female, and more experienced and higher paid. The imbalance is sharpest in states like Texas, which has no middle initial and only an age-derived birth year. The one reversal is Iowa, where linked teachers are *younger* and less experienced; Iowa matches on an age-derived birth year, and older teachers, clustered in fewer birth cohorts, tie against more same-name voters and are dropped more often. Figure A1 shows the distribution of these imbalances across all comparisons, coloured by the field they concern: the excess above the 0.10 threshold is carried overwhelmingly by race and sex, while the grade level an employee instructs is never imbalanced.

Because the fields that predict linkage are observed, we can ask whether this selection distorts the estimates that condition on party, by inverse-probability reweighting. Within each state we model the probability that an employee's party is observed as a function of the recorded demographics, and weight each linked employee by the inverse of that probability, so that the linked sample is reweighted to resemble the full roster on those observables. The weights are stabilized with a state-specific numerator, so that each state's total weight is preserved and the reweighting corrects the demographic composition *within* each state without altering the mix *across* states. We fit this separately per state because the selection differs between state (Texas selects on name commonness, Iowa on age).

We replicate our main result on differential change in district mobility in Table A8 with and without this reweighting. The estimates are essentially identical, and the effective sample size of the weights is 99 percent of the nominal size, so the reweighting is not straining on a few

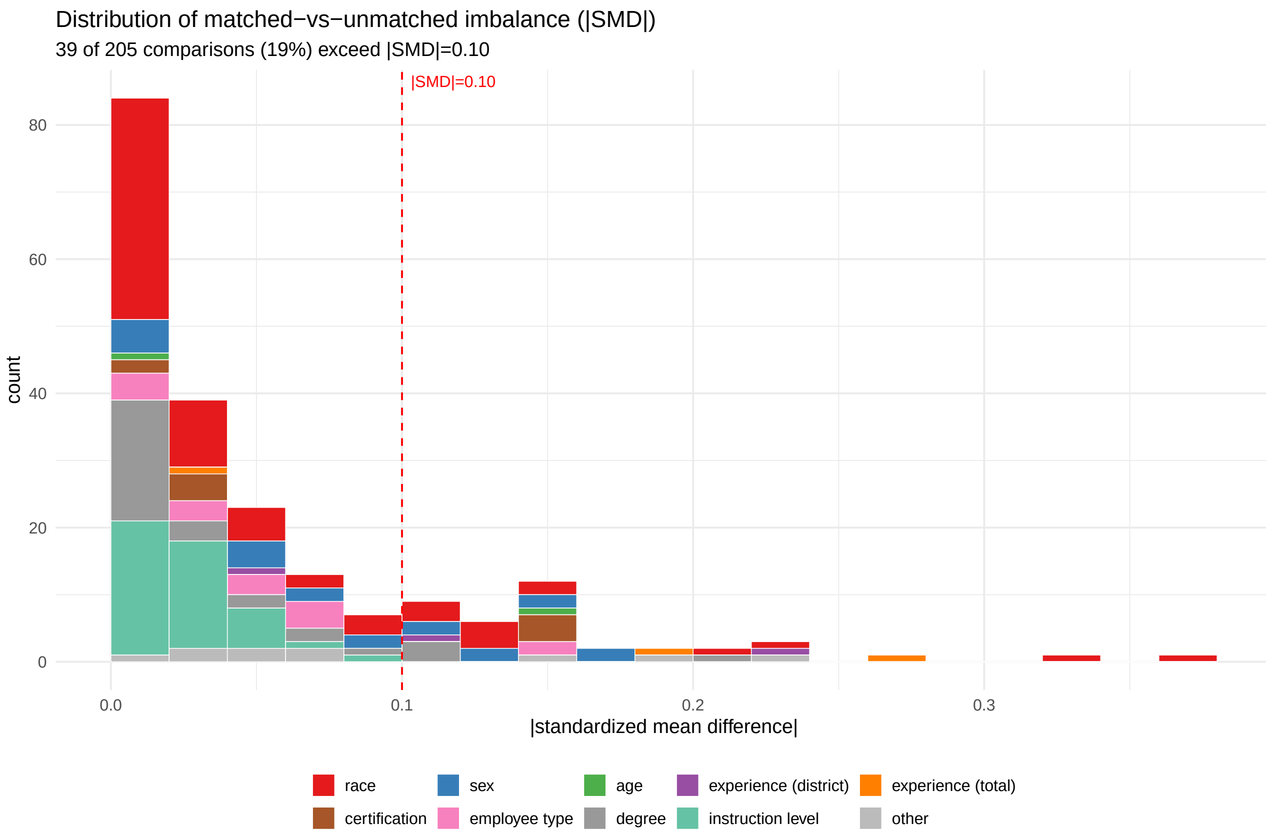

Figure A1. Standardized mean differences between merged and unmerged teachers, across all field-by-state comparisons, coloured by field. The dashed line marks the conventional $|\mathrm{SMD}| = 0.10$ threshold for meaningful imbalance. Two-thirds of comparisons fall below 0.05; the 19 percent that exceed 0.10 are dominated by race and sex, whereas instruction level—the subject or grade taught—never exceeds it.

extreme cases. This indicates that patterns in our merging do not distort our effect of interest, at least *on observables*.

Table A8. Triple-difference in district mobility, with and without IPW reweighting for non-random merging.

	Aligned (R) DiD	Misaligned (D) DiD	Triple $(D-R)$
Unweighted	−0.021 (0.013)	+0.017 (0.008)*	+0.037 (0.013)**
IPW-weighted	−0.022 (0.014)	+0.016 (0.008)*	+0.038 (0.014)**

Note: linked teachers with a Democratic or Republican registration, present at the law's onset, in regular public districts. Fixed effects for matched pair, year, and teacher; standard errors clustered on the treatment-time district. The "Misaligned (D)" column is the sum of the aligned DiD and the triple; its standard error accounts for the covariance between them. IPW weights stabilized per state, effective sample size 99% of nominal. $^{*}\,p<0.05$, $^{**}\,p<0.01$.

D Constructing the Matched Employee Panel

D.1. Objective

The estimates use a matched panel in which every employee who experiences a cultural policy is paired with a comparable employee who does not. We build this panel one treatment source at a time, so the CRT-ban, book-ban, and CATO analyses each have their own matched sample and control group.

D.2. The panel

The unit of analysis for the matching is the employee-year. We reach it in two steps, because employees can work in more than one district or more than one position at a time (although this only applies to 0.5%-0.75% of employees).

We first collapse each state's cleaned roster to the employee-district-year, keeping district in the key so that a employee employed by two districts remains two rows. Within an employee-district-year we take the union of the position, subject, and grade-level indicators, the union of the school-level turnover indicators, the maximum of salary (which is repeated across a employee's position rows), and the first non-missing value for demographics, and we recompute the turnover outcome from its components. Employee identifiers are prefixed by state so that they are unique across the twelve states.

We then create an employee-year panel to use for matching. The indicators and turnover outcomes are unioned across an employee's districts, so an employee counts as treated if any of their districts is treated, and salary is summed, since an employee working in two districts earns from both. Demographics, experience, party, and the remaining covariates are carried forward by the first non-missing value.

The resulting panel covers twelve states and spans 12,455,522 employee-district-years, which roll up to 12,346,489 employee-years for 2,906,566 distinct employees. It is built once and shared across all treatment sources.

D.3. Treatment

We separately build treatment panels which report, for each district and year, the number of enacted policies of a given type. A district's onset is the first year in which its count is at least one. This matters because several districts carry many simultaneous policies. Treatment is absorbing, so an employee is under treatment from the first year in which any district employing them is treated. Control employees are never treated. We merge onset at the teacher-district-year level and then aggregate to avoid losing treatment values for employees working in multiple districts.

D.4. Matching

For each treated employee we take the single row at treatment onset, and we require that they were also observed in each of the two preceding years so that every treated unit carries a two-year pre-history. Controls must satisfy the same two-year requirement.

Pairs must agree exactly on the specific year at treatment and on each of the following: the eighteen position and subject indicators, the four grade-level indicators, party, certified or classified status, full or part-time status, a five-way bin of district experience, a state-specific quartile of age, the district's NCES agency type in that year, and the state. Each of these enters as a factor carrying an explicit missing level, so that unknown values match other unknown values and observed values must agree. Where a state records none of a covariate, that covariate is constant within the state and simply does not constrain its matches. Within each exact cell, we break ties on two continuous district covariates (enrollment and per-pupil spending) by nearest-neighbour Mahalanobis distance, one control per treated unit, with replacement. Matching with replacement means a control may serve several treated teachers, and pair membership rather than employee identity defines the unit of the comparison. The outcome plays no part in the match, so any turnover outcome can be estimated on the same pairs.

D.5. The control pool

Controls are drawn from never-treated employees for both CRT and book bans to avoid possible contamination from the other treatment variables. We do not exclude districts ever treated under CATO due to the scale of this dataset—approximately half of our employees are treated by CATO at some point, so purging CATO-affected controls would bias the control pool toward small and rural districts.

D.6. Exposure and alignment

We define an employee as exposed when they hold any role targeted by a treatment policy. Exposure is one when any target indicator is one and zero only when every target indicator is observed and none is one. It is missing whenever a target indicator is missing and none is one, namely whenever exposure can be neither ruled in nor ruled out. This is deliberately conservative and it forces a substantial share of employee-years to missing, but the alternative would code employee as unexposed on the strength of an absent record, which in the states that do not record subject would silently label their English teachers as unexposed. Missing rows drop from the exposure contrast rather than entering it as zeros.

An employee is misaligned when their party opposes the direction of the policy. Employees who are neither Democrats nor Republicans, and whose party is unknown, are missing rather than aligned.

D.7. Covariate balance

We assess balance by comparing each treated employee with their matched control in the year of treatment onset. We evaluate balance by calculating standardized mean differences across each covariate. We report balance both before matching, where the comparison is against all other employees in the panel, and after matching, where it is against the matched controls, using the same treated-group standard deviation in both so that the two are directly comparable. Figure A2 shows balance for CRT.

Every covariate in the exact cell is balanced perfectly by construction, since each treated employee is matched to a control with identical values, and the standardized difference across all of them is zero. Matching on the age quartile and the experience bin also brings the underlying continuous variables into near-perfect balance, from 0.53 to 0.00 for age and from 0.61 to −0.06 for experience. The covariates left out of the exact cell are balanced incidentally, through their correlation with the covariates that are matched. Sex moves from 0.23 to 0.00, racial composition tightens, and holding a master's degree moves from 0.27 to 0.05. Salary improves from 0.53 to 0.17 but keeps a visible gap, consistent with treated employees working in larger and better-paying districts. The exception is district enrollment: per-pupil spending balances well, moving from −0.69 to 0.07, but enrollment moves only from 1.62 to 0.95 and remains our most imbalanced covariate after matching. Treated employees work in districts far larger than their matched controls. The balance is caused by a lack of common support: CRT bans are overwhelmingly adopted by large districts. We therefore also report results with additional controls for enrollment size in Table A11.

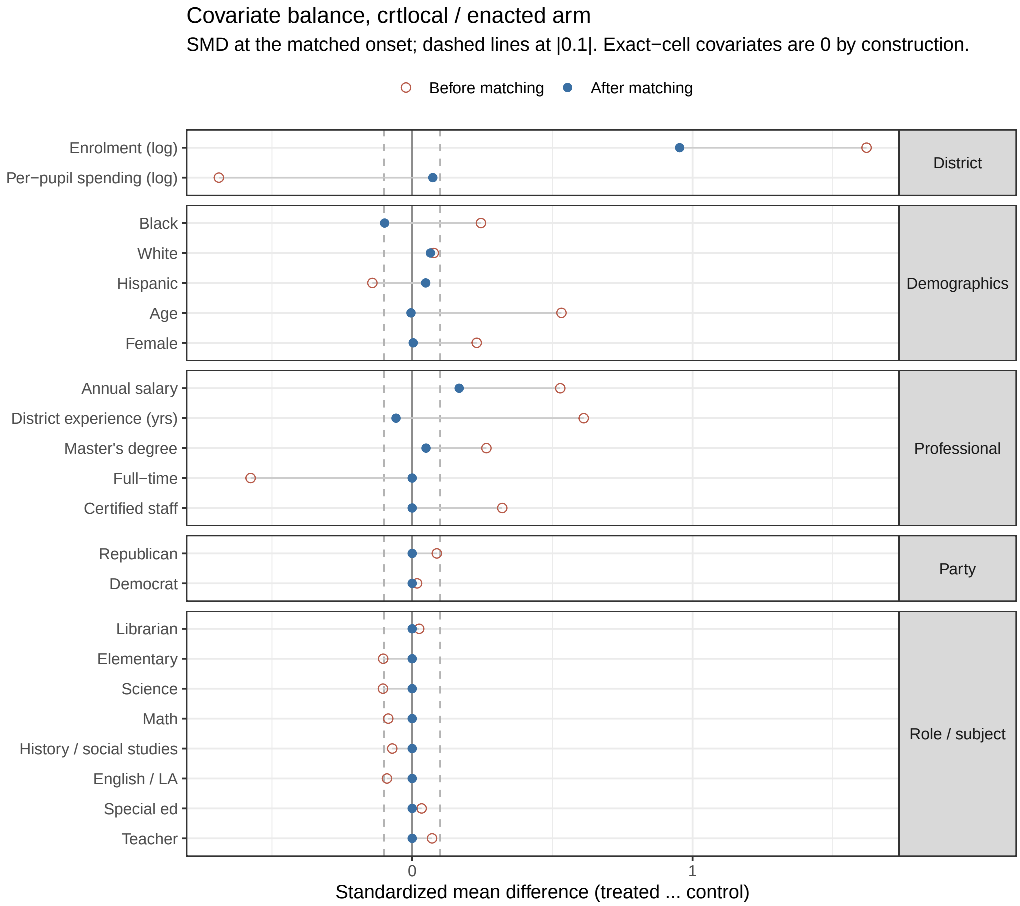

Figure A2. Covariate balance before and after matching, CRT enacted arm. Each point is a standardized mean difference; open circles are before matching and filled circles after, standardized by the treated-group standard deviation. Dashed lines mark |0.1|. Every exact-cell covariate lies on zero.

E Additional Results

E.1. Descriptives

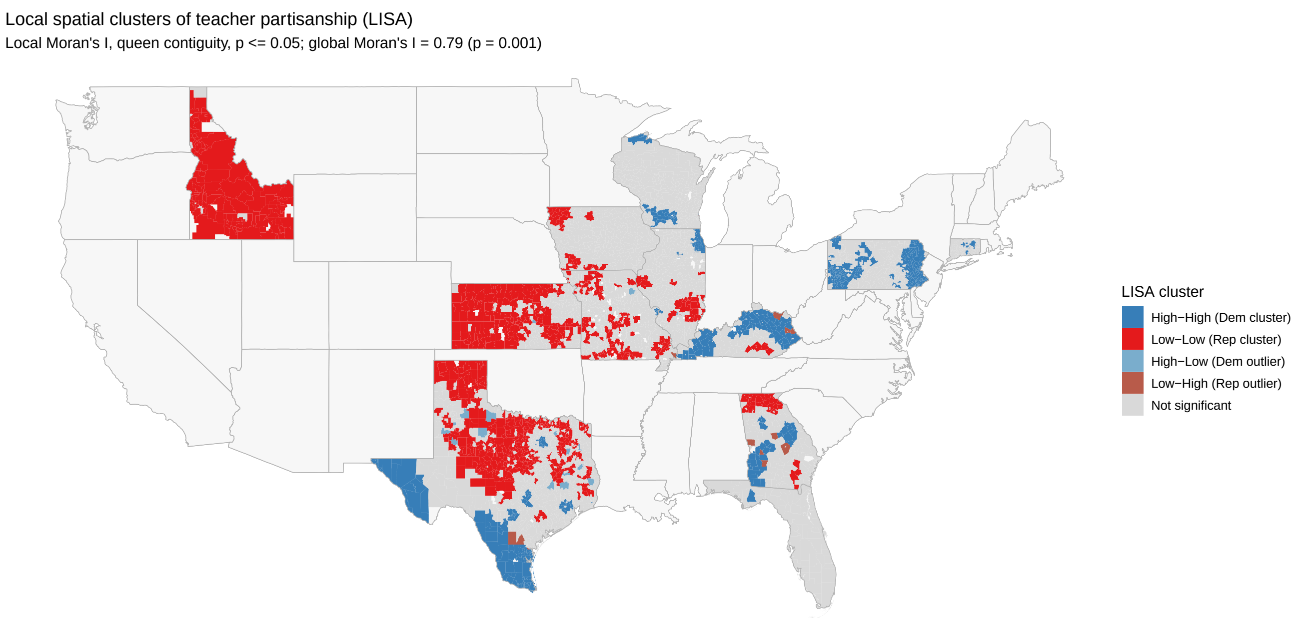

Figure A3. Local spatial clusters of teacher partisanship based on Local Indicators of Spatial Association (LISA) with queen contiguity (neighbors if share a boundary), restricted to clusters significant at p≤ 0.05. High-High clusters (dark blue) are districts with high Democratic teacher shares surrounded by similarly Democratic neighbors; Low-Low clusters (Red) indicate high Republican districts surrounded by similarly Republican neighbors. Outlier categories (High-Low, Low-High) are districts whose partisanship diverges from their immediate neighbors.

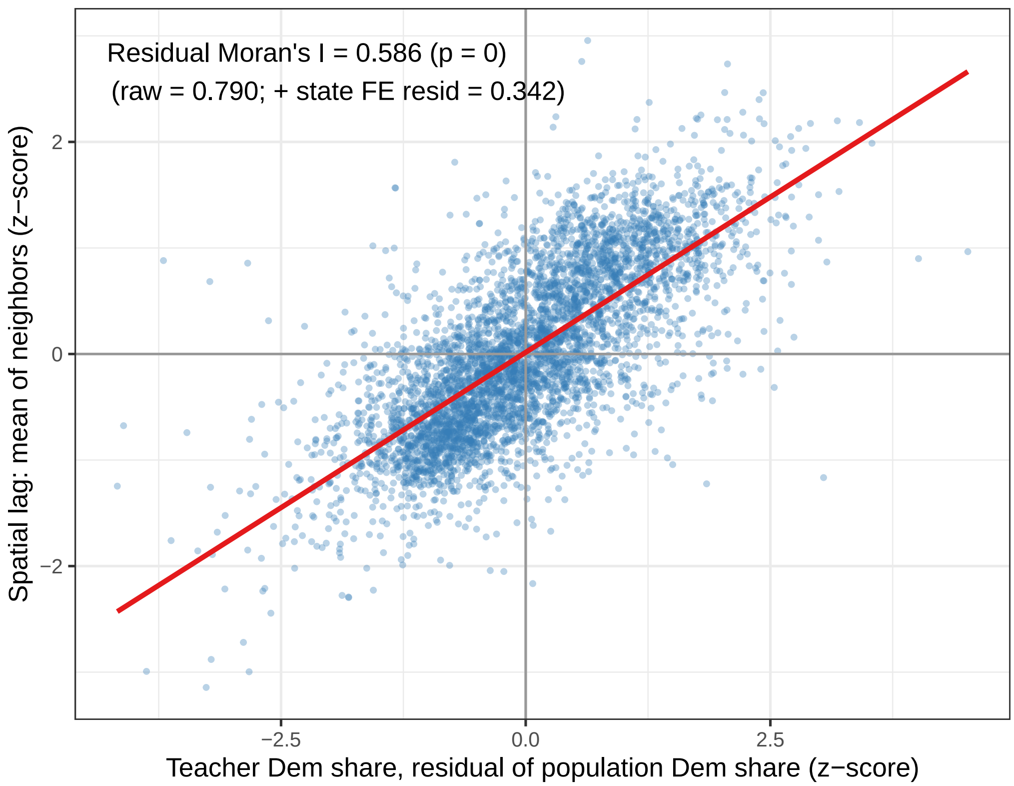

Figure A4. Moran scatterplot of district-level employee partisan composition, residualized with voter population partisanship by school district. Each point is a district; the x-axis shows residualized Democratic share (standardized) and the y-axis shows the mean residualized Democratic share of spatially contiguous neighbors (standardized, queen-contiguity). 12 US states included.

E.2. Turnover Analyses

Table A9. Effect of enacted culturally conflictual policies on last year

Dependent Variable:	Last year on payroll					
Policy	CRT ban		Book ban		CATO	
Spec	DiD	Triple	DiD	Triple	DiD	Triple
Model:	(1)	(2)	(3)	(4)	(5)	(6)
Variables						
post	0.0718***	0.0732***	0.1004***	0.1131***	0.0952***	0.0841***
	(0.0108)	(0.0147)	(0.0061)	(0.0134)	(0.0092)	(0.0064)
treat × post	0.0065	0.0075	-0.0028	-0.0159	-0.0064	-0.0006
	(0.0103)	(0.0135)	(0.0100)	(0.0159)	(0.0084)	(0.0086)
post × misaligned		-0.0031		-0.0254		0.0027
		(0.0122)		(0.0157)		(0.0087)
treat × post × misaligned		-0.0036		0.0249		0.0021
		(0.0130)		(0.0163)		(0.0097)
Fixed-effects						
pair_id	Yes	Yes	Yes	Yes	Yes	Yes
year	Yes	Yes	Yes	Yes	Yes	Yes
teacher_id	Yes	Yes	Yes	Yes	Yes	Yes
Fit statistics						
Observations	1,406,876	971,598	1,391,115	898,413	1,254,766	765,926
R^2	0.24441	0.24207	0.27436	0.27258	0.27004	0.26329
Pre-treat. control mean	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000
Num. pairs	120,806	82,705	132,711	85,376	120,648	73,988

Clustered standard-errors at the district-at-treatment-onset in parentheses
*Signif. Codes: ***: 0.01, **: 0.05, *: 0.1*

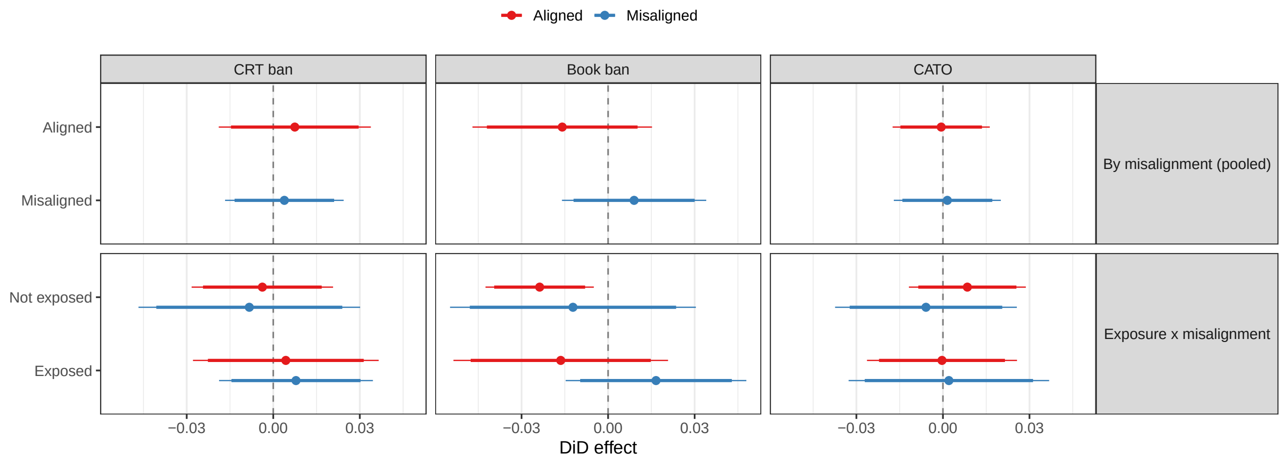

Figure A5. Marginal DiD on leaving the profession, by subgroups and treatment measure, with 90% and 95% confidence intervals.

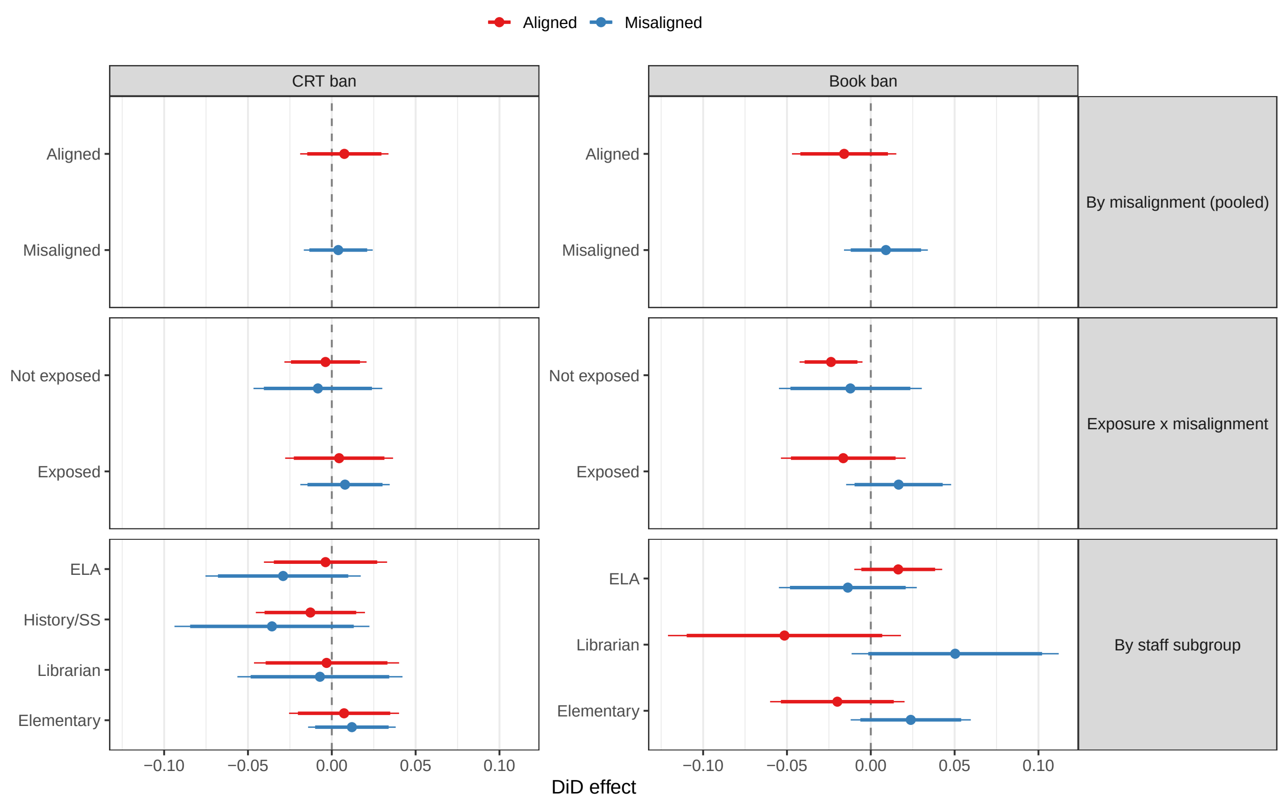

Figure A6. Marginal DiD on leaving the profession, by detailed subgroups and treatment measure, with 90% and 95% confidence intervals.

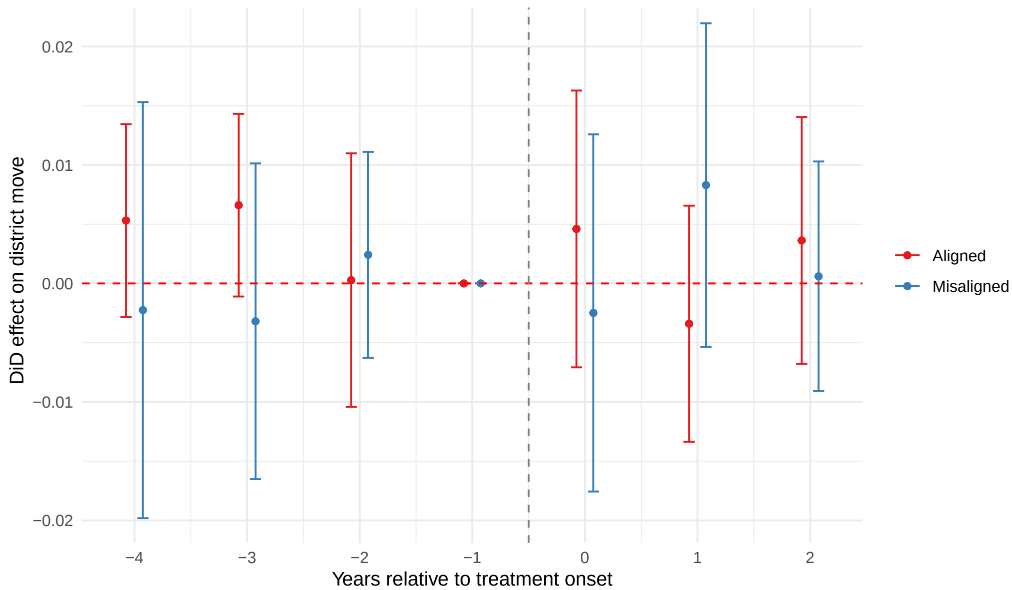

Figure A7. Event study estimates of book bans on district move, with 95% confidence intervals.

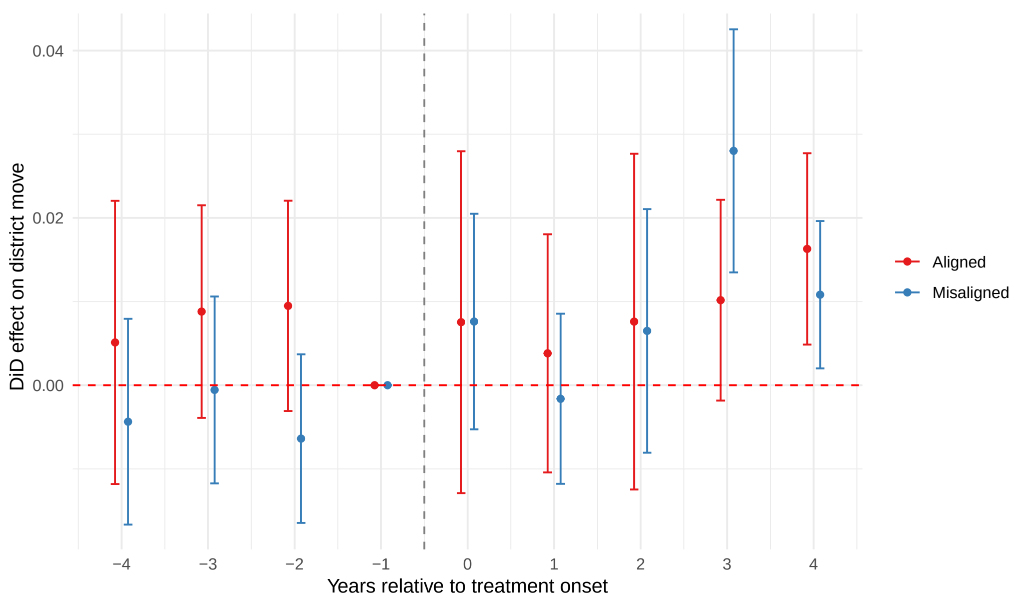

Figure A8. Event study estimates of cultural policies (Cato) on district move, with 95% confidence intervals.

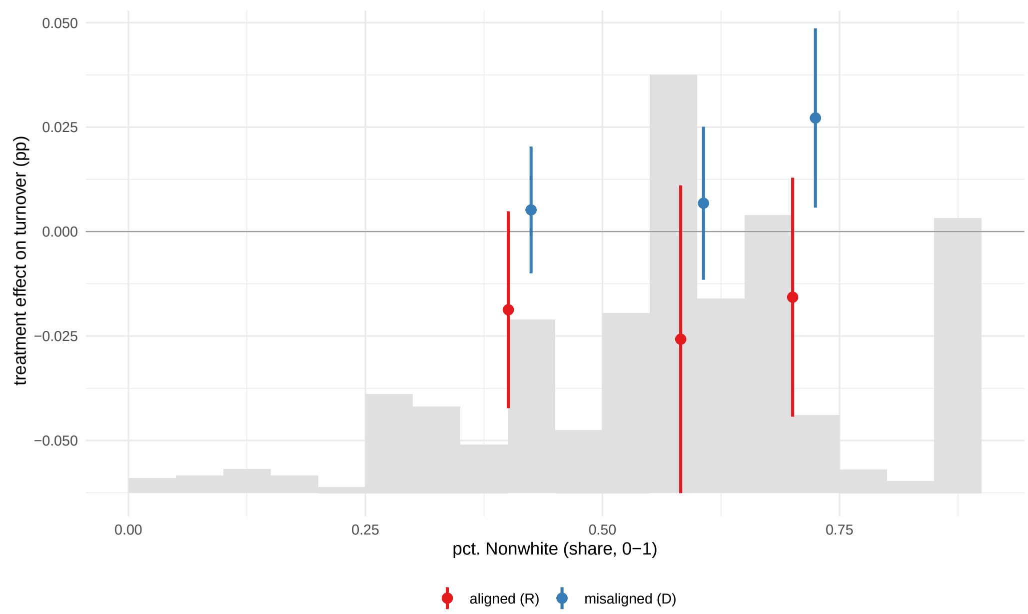

Figure A9. Alignment-specific treatment effect by share of non-White students on district move. Moderator distribution shown in grey.

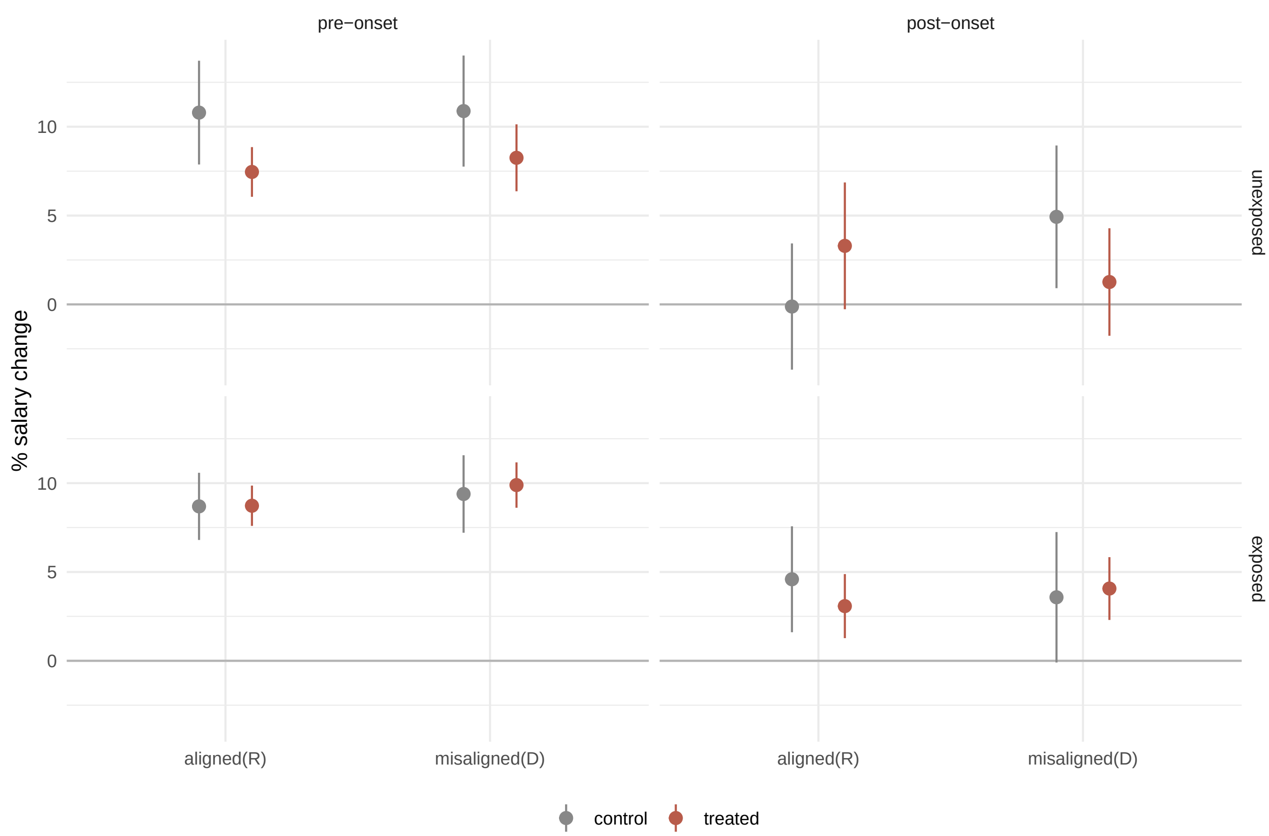

Figure A10. Average percent salary change through a district move, by exposure and alignment, with 95% confidence intervals.

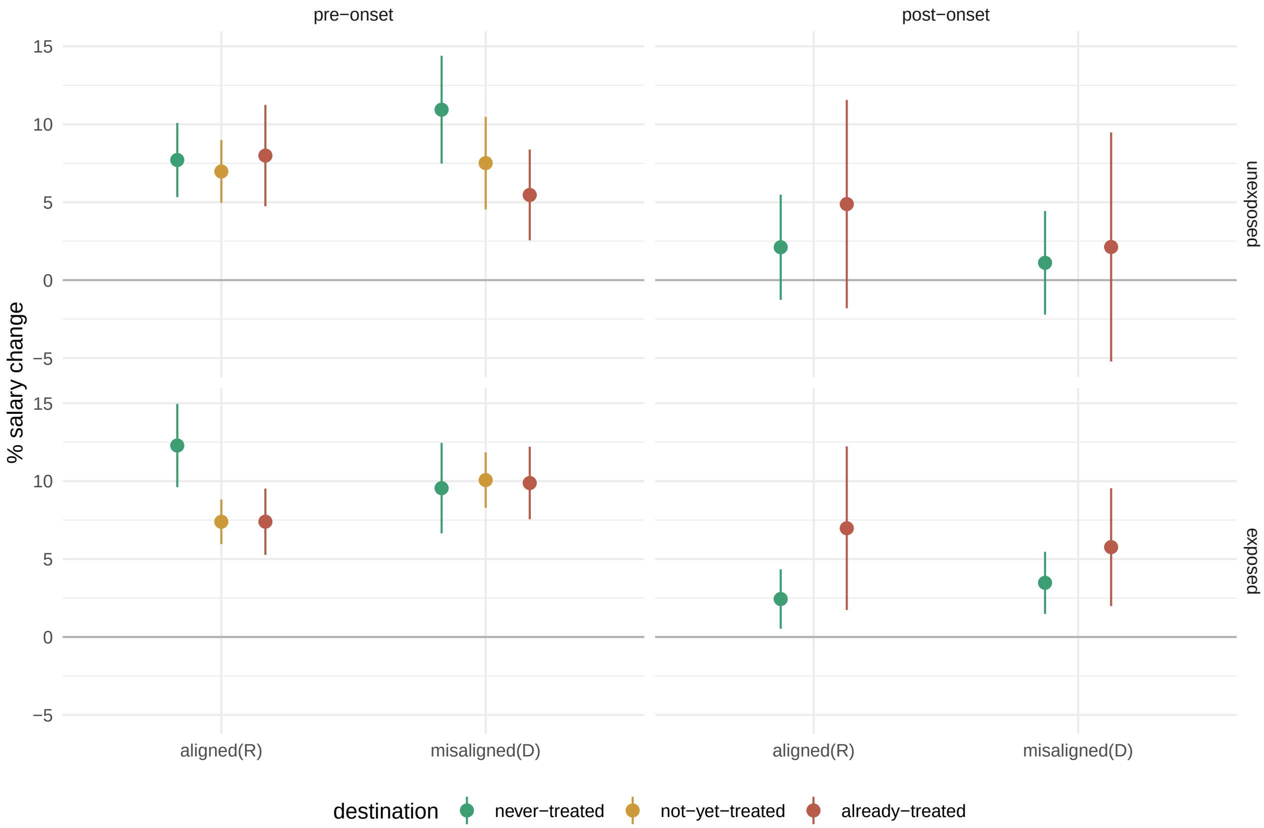

Figure A11. Average percent salary change by destination type for treated movers across districts, with 95% confidence intervals

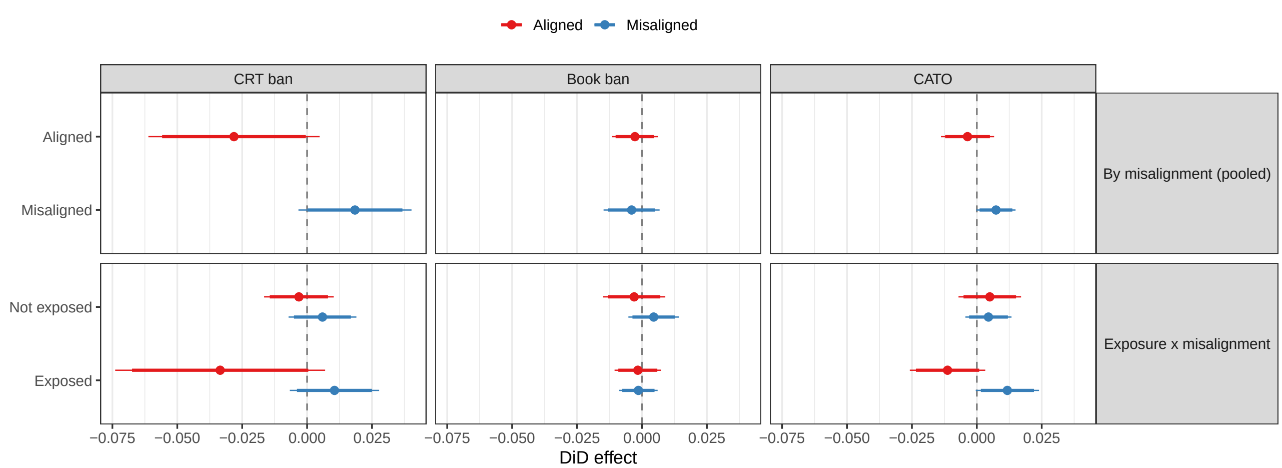

Figure A12. Marginal DiD on district move for teachers only, across subgroups and treatment measures, with 90% and 95% confidence intervals

Table A10. Misalignment triple robustness: CRT, district move

Dependent Variable:	District move			
	(1) Base	(2) +DistrictxPost	(3) +SubjectxPost	(4) +PartyxPost
Model:	(1)	(2)	(3)	(4)
Variables				
post	0.0193**			
	(0.0088)			
treat × post	-0.0207*		-0.0208*	-0.0196*
	(0.0126)		(0.0120)	(0.0104)
post × misaligned	-0.0342***	-0.0379***	-0.0348***	-0.0961**
	(0.0127)	(0.0130)	(0.0132)	(0.0430)
treat × post × misaligned	0.0374***	0.0357***	0.0375***	0.0361***
	(0.0132)	(0.0131)	(0.0134)	(0.0118)
Fixed-effects				
pair_id	Yes	Yes	Yes	Yes
year	Yes	Yes	Yes	Yes
teacher_id	Yes	Yes	Yes	Yes
district at treatment × post		Yes		
subject × post			Yes	
party × post				Yes
Fit statistics				
Observations	971,598	971,418	971,598	971,598
R^2	0.22149	0.22616	0.22214	0.22427

Clustered standard-errors at the district-at-treatment-onset in parentheses
*Signif. Codes: ***: 0.01, **: 0.05, *: 0.1*

Table A11. Effect of enacted culturally conflictual policies on district move, controlling for district size (enrollment quintile x post)

Dependent Variable:	District move					
Policy	CRT ban		Book ban		CATO	
Spec	DiD	Triple	DiD	Triple	DiD	Triple
Model:	(1)	(2)	(3)	(4)	(5)	(6)
Variables						
treat × post	-0.0051	-0.0222	-0.0024	-0.0010	0.0024	0.0019
	(0.0075)	(0.0136)	(0.0033)	(0.0039)	(0.0036)	(0.0051)
post × misaligned		-0.0308***		-0.0037		-0.0036
		(0.0107)		(0.0058)		(0.0039)
treat × post × misaligned		0.0341***		-0.0002		0.0061
		(0.0107)		(0.0063)		(0.0049)
Fixed-effects						
pair_id	Yes	Yes	Yes	Yes	Yes	Yes
year	Yes	Yes	Yes	Yes	Yes	Yes
teacher_id	Yes	Yes	Yes	Yes	Yes	Yes
enrollment quintile × post	Yes	Yes	Yes	Yes	Yes	Yes
Fit statistics						
Observations	1,406,876	971,598	1,391,115	898,413	1,254,766	765,926
R^2	0.22495	0.22174	0.25724	0.25650	0.25182	0.25071
Pre-treat. control mean	0.0304	0.0323	0.0203	0.0205	0.0255	0.0262
Num. pairs	120,806	82,705	132,711	85,376	120,648	73,988

Clustered standard-errors at the district-at-treatment in parentheses
*Signif. Codes: ***: 0.01, **: 0.05, *: 0.1*

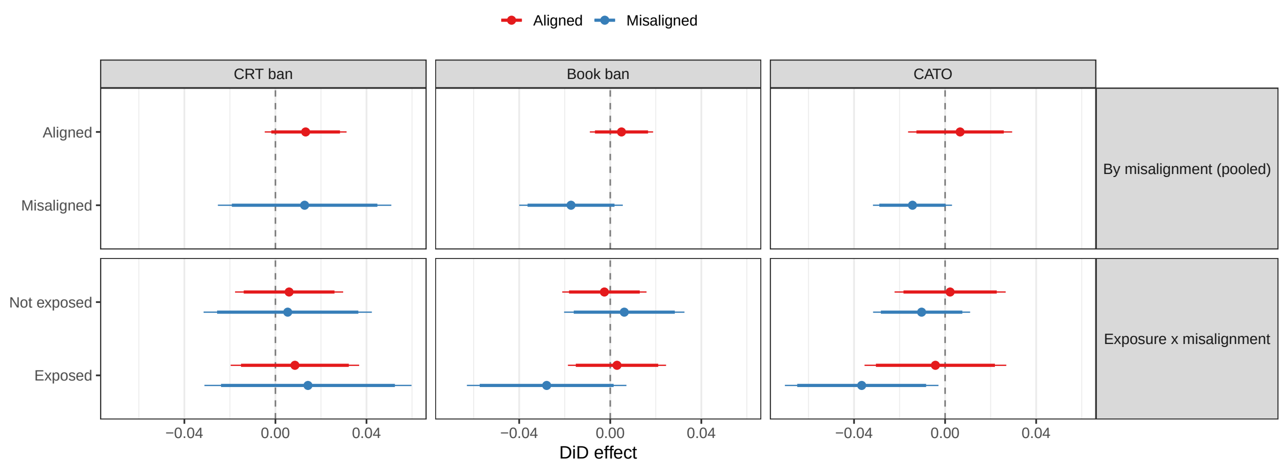

Figure A13. Marginal DiD on school move, across subgroups and treatment measures, with 90% and 95% confidence intervals

F Details on Qualitative Sampling and Teacher Interviews

We complement our empirical analyses with qualitative interviews with 10 current and former K-12 public school employees. This component of the paper was approved by Stanford IRB-FY2026-87527.

F.1. Residual-Based Case Selection Strategy

To select school districts from which to contact employees for interviews, we employ a common mixed-methods approach that uses the regression model itself to identify districts where the model fits well ("typical" cases that minimize the size of the residuals – namely, the distance between the observed and predicted values of the outcome) versus cases where observed outcomes depart markedly from the model predictions ("deviant" cases that maximize the size of the residuals) (Nalewajko, 2025; Seawright, 2016). Typical cases are useful for process tracing the causal mechanisms, allowing us to observe whether the theorized process – shifts in employees' alignment with the institution's official position and constraints on their discretion – are operating in the cases where the model fits well. By contrast, deviant cases identify districts where turnover diverges substantially from what the model predicts (in either direction), pointing to conditions that are not captured by the main specification and potentially revealing important scope conditions. For example, a deviant-treatment district with turnover that is higher than predicted may point to conditions that amplify the costs we identify, such as enforcement intensity.

To identify the districts for comparison, we use the main specification (Equation 6) with the CRT data and district-switching outcome. We compute district-level residuals that capture the gap between each district's observed post-period turnover and what the model predicts given its treatment status and employee composition. We standardize the residuals and classify all of the districts roughly along two dimensions: typical cases fall within the lowest decile of absolute residuals (clustering close to the "0" residual value) and deviant cases fall within the top decile (further away from "0"). Crossing the treatment status with this classification yields four bins of districts from which we select our cases: typical treatment, deviant treatment, typical control, and deviant control. One feature of the data is worth noting: no treatment district formally clears the top-decile global deviance threshold, as all treatment-side residuals fall within roughly ± 0.30 standard deviations of zero; we therefore designate these treated districts as deviant based on the large absolute residuals among the treated districts.

We select 19 districts from which to sample employees for interview outreach: six typical treatment, six deviant treatment, four deviant control, and three typical control. Beyond residuals, we also took account of several other factors in our selection process. Following similar insights of informal blocking from (Nalewajko, 2025), we aimed for the school districts to satisfy the following conditions:

1. Are "typical" or "deviant" – namely, they have low or high residuals.
2. Are different from one another according to district enrollment size, with a preference placed on medium and larger-sized districts for greater likelihood of response.
3. Are different from one another according to sources of conflict across our datasets, ranging from one source (anti-CRT) to three (anti-CRT, book bans, and CATO policy).
4. Have reasonable contact feasibility (e.g. standardized professional email formats or school district email directories, particularly for employees who have not left the district or profession). This has a bias toward reaching current staff that we partially offset with L2 voter-file phone records, which are more effective for reaching employees who have already left.

We suppress information about the selected districts to maximally maintain respondent confidentiality.

F.2. Employee Sampling Strategy Within Selected Case Districts

For each of the selected districts, we identify employees for potential outreach using the full state-level rosters files, but we privilege the employees who are in the matched panel. This approach captures all employees who work in the district at the relevant year. For treatment districts, we use the earliest year the CRT local policy took effect; for control districts, we use the onset year of the matched treatment pair. We then cross-reference the roster against the matched panel to identify which employees also appear in the analytical sample. Panel-matched employees receive a cell designation based on the two theoretical dimensions: subject-area exposure to the policy (exposed vs. unexposed) and voter-ideology alignment (misaligned vs. aligned), producing four cells.

To identify the employees that we contacted for interviews from this pool of 66,767 total candidates, we merge contact-information (phone numbers) from L2, which are important for outreach to employees who have left the school district, retired, or otherwise exited, and district staff directories or indicators of standardized district email formats. Following this merge, we draw a stratified random sample for outreach and target cell sizes based on potential response rates and projected numbers of interviews based on those response rates. For treatment districts, we select an equal number of employees per cell, targeting about ten per cell, to ensure balanced coverage of the exposed-by-ideology matrix across all four cells. For control districts, we give priority to employees who appear in the matched panel, then supplement with a random draw from the non-panel roster to reach fifteen contacts per district. Using information on the projected patterns of district emails (e.g, firstname.lastname@domain, firstinitiallastname@domain, lastnamefirstinitial@domain), we used Claude Sonnet 4.6 to construct the teachers' emails using their first and last name and sent out emails in batches.

F.3. Additional Mixed Outreach Approaches

Our outreach efforts in mid July and August 2026 faced a number of emails that bounced back (251 out of 606 emails bounced), a 41.4% bounce rate overall from the constructed emails. Some of these are expected given that 76 employees have positive turnover in the administrative records.). We also faced high rates of non-response. We expect that the non-reply rate is shaped by four factors: employees being on summer break and not checking their emails until the school year begins, emails ending up in spam folders, disinterest in participation, or concerns over participation due to the restrictive atmosphere. Regarding the latter reason, we received a response from an educator (paraphrased) who mentioned they are not allowed to do interviews to speak about how recent policies have influenced teachers' work in the district, but that it has been a challenging past few years where they are.

Given these different constraints, we combine our principled sampling approach with several other strategies: resending the emails to our original sample closer to the start of the school year, outreach to teachers' union representatives within districts, and snowballing sampling among employees who do respond to our emails.

F.4. Interview Structure and Sample Questions

Interviews with subjects are semi-structured, beginning with the informed consent process and organized into a common opening, a substantive block loosely tailored to the district case type, and a closing.

Opening: All respondents are asked to describe their teaching career in their own words: what motivated them to pursue teaching, how long they have taught, where and what types of schools they have worked in, what subjects they have taught, and what they are doing now. Respondents who have left the classroom are also asked about their current role.

Substantive Block: The interview pivots to state and local educational policy with an open prompt, asking whether anything comes to mind regarding political debates or new policies – over instructional content, teaching techniques, or classroom speech – that have affected the respondent's work in recent years. Follow-up questions ask how the interviewee first heard about the incident(s); whether the debates or policies disrupted the employees' own day-to-day work or that of colleagues; and what themes come to mind in connection with these debates (e.g., effects on job satisfaction, autonomy, morale, sense of security, or identity as a public school teacher, and whether these effects have improved or worsened, and for whom). A broader question then asks how the employee feels about the ways national politics can affect the neutrality of institutions like schools, politicize them, and what this means for teachers' sense of mission and for their students.

The interview then turns to employment decisions: whether they or anyone they know has considered leaving teaching entirely or changing school districts; for those who left, what drove the decision; for those who stayed, whether staying was an active choice (e.g., higher pay, a supportive principal, colleagues, or union, independent of the broader district climate) or reflected personal or professional constraints (geographic ties, a difficult job market, tenure considerations, proximity to retirement).

Closing questions ask whether there is anything about the respondent's experience as an educator, in relation to these topics or more broadly, that the interview did not cover, and whether they colleagues at their school or elsewhere who might be open to speaking with the research team.

G Additional Qualitative Findings

Table A12. Participant Characteristics, Completed Interviews

Role	Case Type	Cell	Any Turnover	Approx. Tenure	Sampling Frame
science teacher	Typical Treat	A	No	~35 years	Sampled state
math teacher	Typical Treat	C	No	~25 years	Sampled state
ESE/reading teacher	Typical Treat	A	No	~30 years	Sampled state
elementary teacher	Typical Treat	B	No	~6-9 years	Sampled state
elementary teacher	Deviant Treat	A	Yes (self-reported)	~7 years	Sampled state
world-language teacher	Typical Control	D	No	~23 years	Sampled state
administrator (former)	Deviant Treat	A	Yes (self-reported)	~9 years	Sampled state
librarian	Outside Sampling	A	No	~16 years	Snowball referral
administrator	Typical Treat	B	No	~9 years	Sampled state
science teacher (private school)	Outside Sampling	Unclassified	No	~20+ years	Snowball referral

Notes: District names and exact tenure are generalized to protect participant anonymity. *Sampling Frame* flags respondents reached outside the 19-district, four-state design (e.g. via snowball referral).

Table A13. Additional Quotes: Theorized Channels. A = Exposed + Misaligned; B = Exposed + Aligned; C = Not Exposed + Misaligned; D = Not Exposed + Aligned.

Role	Construct	Quote
Cell A		
science teacher	Alignment	"I think a lot of teachers feel like they're expected to be a bot"
science teacher	Discretion	"we used to always do the sex education, and in middle school, we don't do it anymore"
ESE/reading teacher	Alignment	"teaching was not considered a professional degree anymore"
ESE/reading teacher	Alignment	"there's less trust that we're the ones who know what we're doing"
ESE/reading teacher	Discretion	"the shelves are 50% of what they used to be"
librarian	Alignment	"I think the plan is to drive us all out and replace us all with 22-year-old [ideologically aligned recent graduates]"
librarian	Discretion	"it is all cut and paste...no personalization. Teachers can't teach what they want at all"
Cell B		
elementary teacher	Discretion	"a colleague in a same-sex marriage described being unable to display a family photo in her classroom, unlike other staff"
administrator	Alignment	"that language (don't say gay) was not in the legislation anywhere. That was a sound bite"
Cell C		
math teacher	Discretion	"we adopted a (math) textbook, and then it came back and said we couldn't adopt that textbook because it was too woke"
math teacher	Discretion	"I have been much more careful about what I post on my personal social media"
Cell D		
world-language teacher	Alignment	"I have been frustrated by the ways that we are coached on how to think about our own culture...pressed to be sensitive to every foreign culture, and to suppress our own"
world-language teacher	Alignment	"I vote against it (CRT) every time I get a chance. But I still signed contracts, and here I am."
Outside cell framework (private school / non-panel)		
science teacher	Alignment	"other people coming from some other field...saying that they can be a teacher...that makes me feel like my job is cheapened"
science teacher	Discretion	"if I was in the public school, would I be able to do this? I really doubt it"

Table A14. Other Challenges Teachers Raised, Beyond Culturally Conflictual Policy. These quotes illustrate prominent challenges respondents raised unprompted that are distinct from the alignment and discretion channels our framework theorizes. We report them for transparency and to situate our findings within the broader landscape of pressures on the teaching profession, not as evidence for or against our framework.

Role	**Case Type**	**Quote**
Behavior/discipline		
elementary teacher	Typical Treatment	"one of the big driving factors that has a lot of teachers leaving is the behavior and the lack of consistency with behavior...But I definitely think that behavior is a major driving factor for a lot of teachers leaving"
elementary teacher	Typical Treatment	"out of an entire public elementary school that was known for throwing desks and...having these crazy outbursts, and now there's one in every classroom"
administrator (former)	Deviant Treatment	"the other thing that's been a looming issue is, of course, behavior...you do have a lot of behavior concerns, especially whether it be inner city or a small rural setting. And sometimes those special education rights that are given to students...schools are really apprehensive of essentially removing a student that might be impeding other students' ability to learn, because they are much more concerned about dealing with any legal matter...so I am seeing a lot of times that behavior or the inability to manage behavior within the building from an admin standpoint has definitely impeded teachers' morale"
librarian	Outside Sampling Frame	"part of the reason I left the last elementary school was because I would get punched on a regular basis going into certain classrooms. And the fact that they were special ed classrooms doesn't really matter...we were being pretty regularly assaulted"
science teacher (private school)	Outside Sampling Frame	"they called me [a string of profanity and insults]...I would call home, and [the response was] we don't consider goddamn a cuss word...when I would go to administration, and I would send the kid out of the room, the kid would be sent back to the room, saying that we need him in the classroom to learn, not out here in the office or out in the hall. I really did not know how to handle that. It's like, I'm...I have no control over my class with this type of stuff going on"
Pandemic/COVID		
world-language teacher	Typical Control	"probably when COVID happened and schools all went online, that was the time when I saw people stop teaching. And that was the time that I seriously considered not teaching. It's not because I don't like teaching, it's because there was no teaching happening...there was zero teaching happening...and the policies that the counties put in place were ridiculous. And they created a lot of problems...it taught a lot of bad habits for teachers and students, and it's been hard to come back from"

continued on next page

Table A14 (continued)

Role	Case Type	Quote
elementary teacher	Typical Treatment	"we were half online, half in person. So...I was teaching 5th grade science and social studies that year, and I would have...anywhere from 5 to 10 kids online, and then...10 to 15 in the classroom at the same time, and I had 3 different classes"
science teacher	Typical Treatment	"that changed, I would say, 6 years ago with COVID, because parents were ignorant of what was going on in the school system, and ignorance can be bliss. So, we started getting parents literally getting in our faces when we had to do...the home teaching, and a parent would literally be in the classroom, and they're like, what are you teaching my child?"
Technology/devices/AI		
science teacher (private school)	Outside Sampling Frame	"I think technology is starting to ruin us. I'm scared half to death with AI. Can kids really do anything on their own and think?...Something needs to be done about AI. We're taking away our kids' creativity and their thinking. I can't tell you when I do science fair how many of them just look, and they just put the whole project, their abstract, everything, in AI"
ESE/reading teacher	Typical Treatment	"there is a bit of a movement to get away from so much digital content, which is a very good thing, because we really have seen that kids being on screen like 90% of their waking hours is absolutely having a detrimental effect on them. Whether it's school screen or home screen, it doesn't matter...they just don't know what to do with the real world anymore and younger and younger...even kids with no disabilities or anything, just, like, 10 and 11-year-olds, can't maintain eye contact"
world-language teacher	Typical Control	"it's about time we swing back to merit-based education...that we completely eradicate and reject the no right answers, just good tries...that we get devices and technology out of our classrooms, because it doesn't help, it hinders. Our kids can't think, they can't do simple things, they don't reason well. High school kids function, many of them, at a level of second or third grade"
elementary teacher	Typical Treatment	"I pull from my own experience in the sense that I was used to seeing that kind of stuff growing up, and so I feel like this generation has really missed out...even with the whole tablet situation, everything, right, this generation already lacks basic common decency and interaction, right, because of the tablets"

H Illustrative DistrictView Quotes

Table A15. Illustrative Teacher Public Comments on Culturally Contentious Policies. Each row is one public comment by a self-identified current or former district employee, identified using transcripts from school board meetings in *DistrictView* (Holman et al., 2025). Quotes are drawn from districts identified as treated in our CRT analysis. The links in the Recording column open the meeting video at the moment the comment begins.

Public Comment	District, State	Date	Recording
[A 19-year history teacher, on a proposed policy for teaching controversial issues:] "For most of my career I have felt myself to be a valued and respected member of the community... Lately, though, I have not felt so valued and respected. The proposed policy on teaching controversial topics would name teachers as the enemy, the boogeyman, the villain. The proposed policy declares that my colleagues and I are not to be trusted... I am not a Marxist. I am not a fascist. I am not a groomer. I don't teach critical race theory. No teacher I know of in this county fits any of these descriptions... Teachers will be nervous to teach any issue that might be considered controversial and will find ways to avoid the topics. I see it already happening among my colleagues, with teachers second-guessing what would be normal lessons and student-driven class discussions—and you haven't even voted on the proposal."	Isle of Wight County, Virginia	Feb. 9, 2023	`XC0lNLgZ6iA` 1:08:07
[A 23-year teacher, speaking against an anti-CRT resolution on that night's agenda:] "These two resolutions are very concerning to me... I feel upset for a number of reasons. This [anti-CRT] resolution stifles discussion and may result in teachers feeling unable to answer uncomfortable questions from students or to have honest conversations and discussions about how race affects our lived experiences... If what we want to do is ensure a safe inclusive environment for all our students, vote against this unnecessary and misguided resolution. It works against the very idea of anti-racism, and it may make teachers feel unable to engage in honest dialogue about important and relevant issues to students and their lives, like racism."	Temecula Valley USD, California	Dec. 13, 2022	`eosZ1Tf6Fl0` 3:51:45

continued on next page

Table A15 (continued)

Public Comment	**District, State**	**Date**	**Recording**
[A fifth-grade teacher who served on the district's social studies curriculum pilot team:] "I would like you to think back to just about two months ago when you brought in self-proclaimed experts to explain your vision on CRT. You considered these people experts... One of the attendees asked, would you consider us experts in our classroom, and [his] response was yes... I am asking now that you consider our expertise when considering the social studies adoption... We don't want to bring contention into our classroom or put ourselves in front of the firing squad, per se... Show your trust in our expertise like you showed your trust in the self-proclaimed experts a few months ago."	Temecula Valley USD, California	May 16, 2023	ABcKfZu7_pU 1:02:49
"[Special education teachers] are tired and feeling burnt out. Please consider ways to help them serve our students with some of the largest needs. Please stop talking about politically driven topics such as critical race theory—which the board chair referred to as an issue that receives only an occasional email—and start talking about what is important: finding ways to retain and recruit much needed educators."	Berkeley 01, South Carolina	Mar. 7, 2023	eWE3PCzZ7gM 1:38:22
[*A former teacher speaking in favor of a proposed CRT ban*]: "You work for us, and Mr. Superintendent, we pay for your salary. Now like I say, I don't know where you people come from or why, who you think you are that you can accuse my grandson of being a racist... my grandson is not a racist, and I will tell you from my family, myself, and I would say the overwhelming majority of people at this meeting, that I say no to the 1619 Project, I say no to the slanderous portrayal of white privilege, I say no to this screwing around of the definition of equality and equal and equity, and I say no to the critical race theory. I say yes to banning—your Policy #832."	Kutztown Area SD, Pennsylvania	May 17, 2021	BGPSSeUj91O 3:00:33

References

Enamorado, T., Fifield, B., and Imai, K. (2019). Using a probabilistic model to assist merging of large-scale administrative records. *American Political Science Review*, 113(2):353–371.

Holman, M., Johnson, R., and Simko, T. (2025). Measuring conflict in local politics. *Urban Affairs Review*, page 10780874251355893.

Nalewajko, K. (2025). Allies of the weak: La Résistance and Jews in the Holocaust. *American Political Science Review*, 119(4):1942–1962.

Seawright, J. (2016). *Multi-method social science: Combining qualitative and quantitative tools*. Cambridge University Press.